

# **CMSC 332**

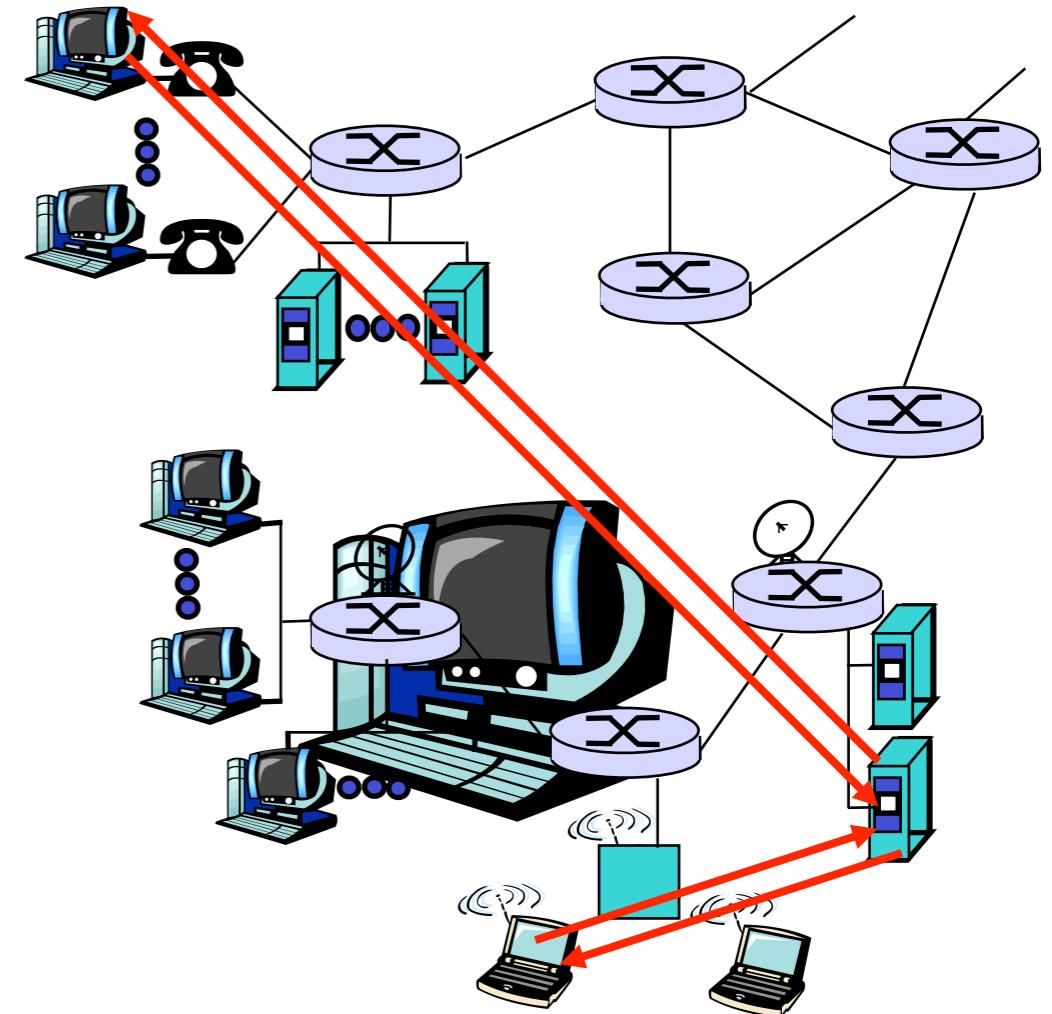
# **Computer Networking**

## **Web and FTP**

**Professor Szajda**

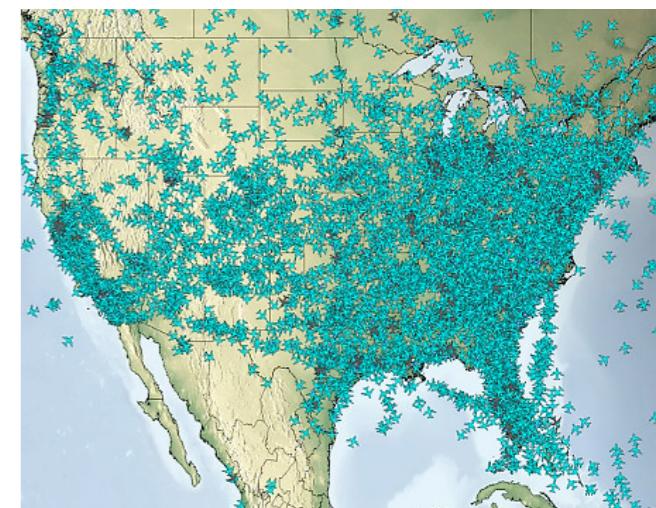
# Review

- In the last slide set, we talked about principles of network applications
  - ▶ End-to-end argument
  - ▶ Network architectures (Client/Server, P2P)
  - ▶ Service requirements



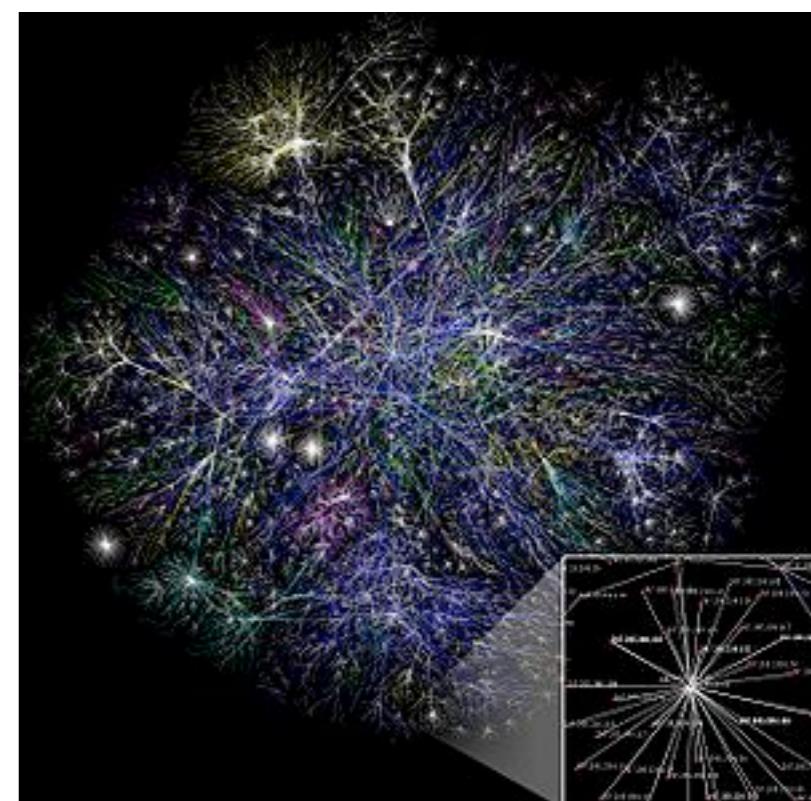
# Aside: Bandwidth-Delay Product

- A student once asked about the “bandwidth-delay product”.
- This is simply the bandwidth of a link multiplied by the propagation delay (in seconds).
  - ▶ It tells us how many bits can fit “in the pipe”.
- Example: If we have a 10Mbps link between here and Berkeley (with a 100ms delay), what is the bandwidth-delay product?
  - ▶  $10\text{Mbps} * 1/10\text{sec} = 1 \text{ Mb}$



# Chapter 2: Application layer

- 2.1 Principles of network applications
  - 2.2 Web and HTTP
  - 2.3 FTP
  - 2.4 Electronic Mail
  - 2.5 DNS
  - 2.6 P2P file sharing
  - 2.7-2.8 Sockets
  - 2.9 Building a webserver



# Web and (Basic) HTTP

# First some jargon

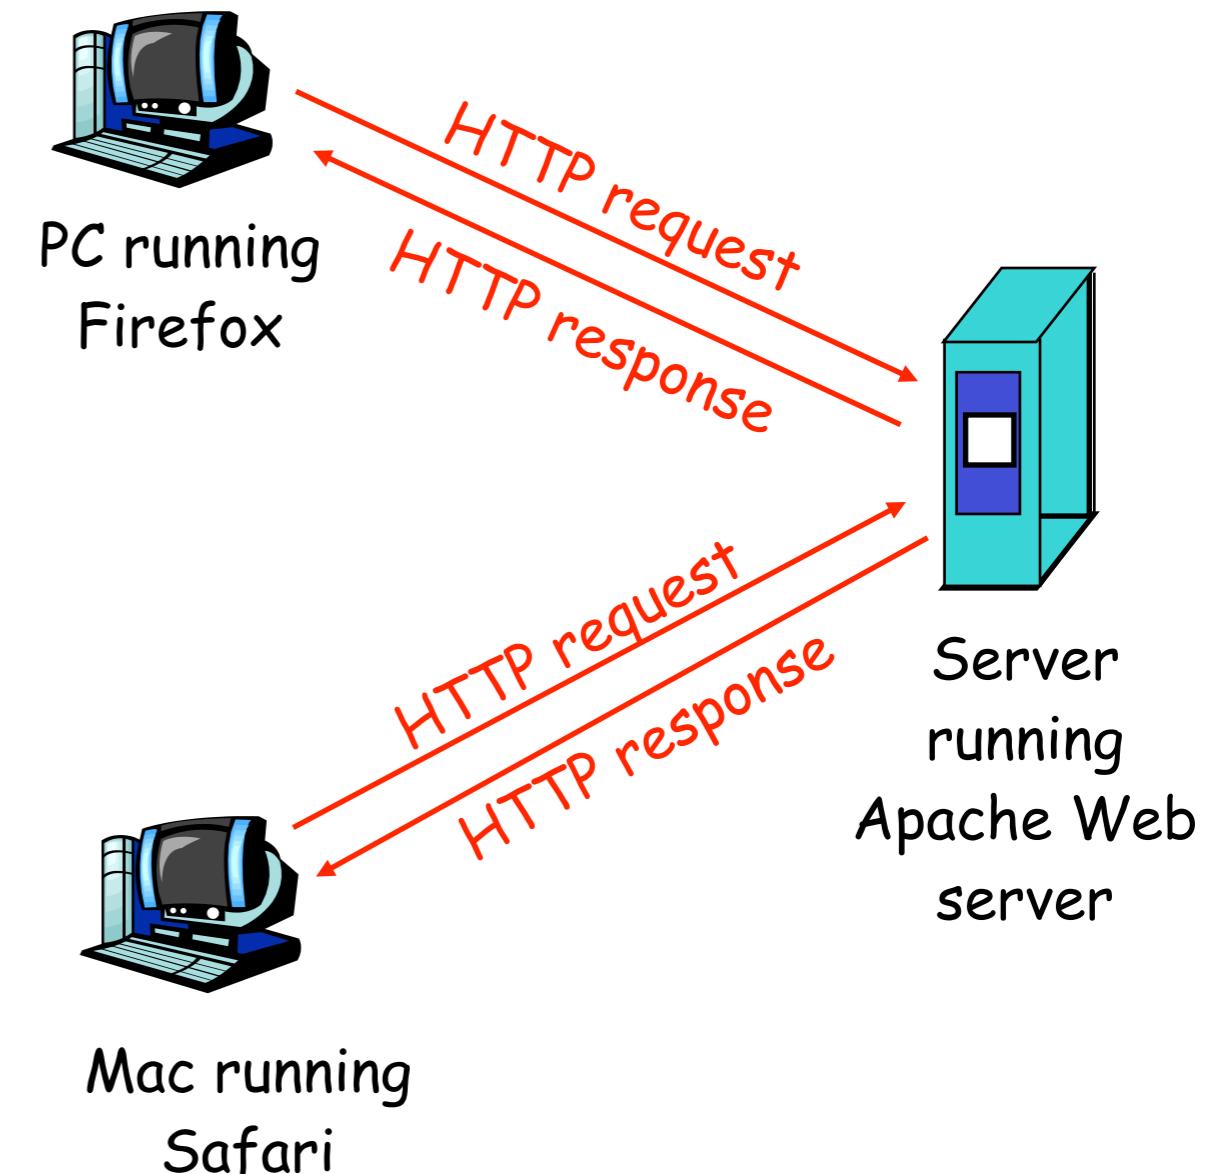
- Web page consists of objects
  - Object can be HTML file, JPEG image, Java applet, audio file,...
  - Web page consists of base HTML-file which includes several referenced objects
  - Each object is addressable by a URL
  - Example URL:

www.someschool.edu/someDept/pic.gif

# HTTP overview

## HTTP: hypertext transfer protocol

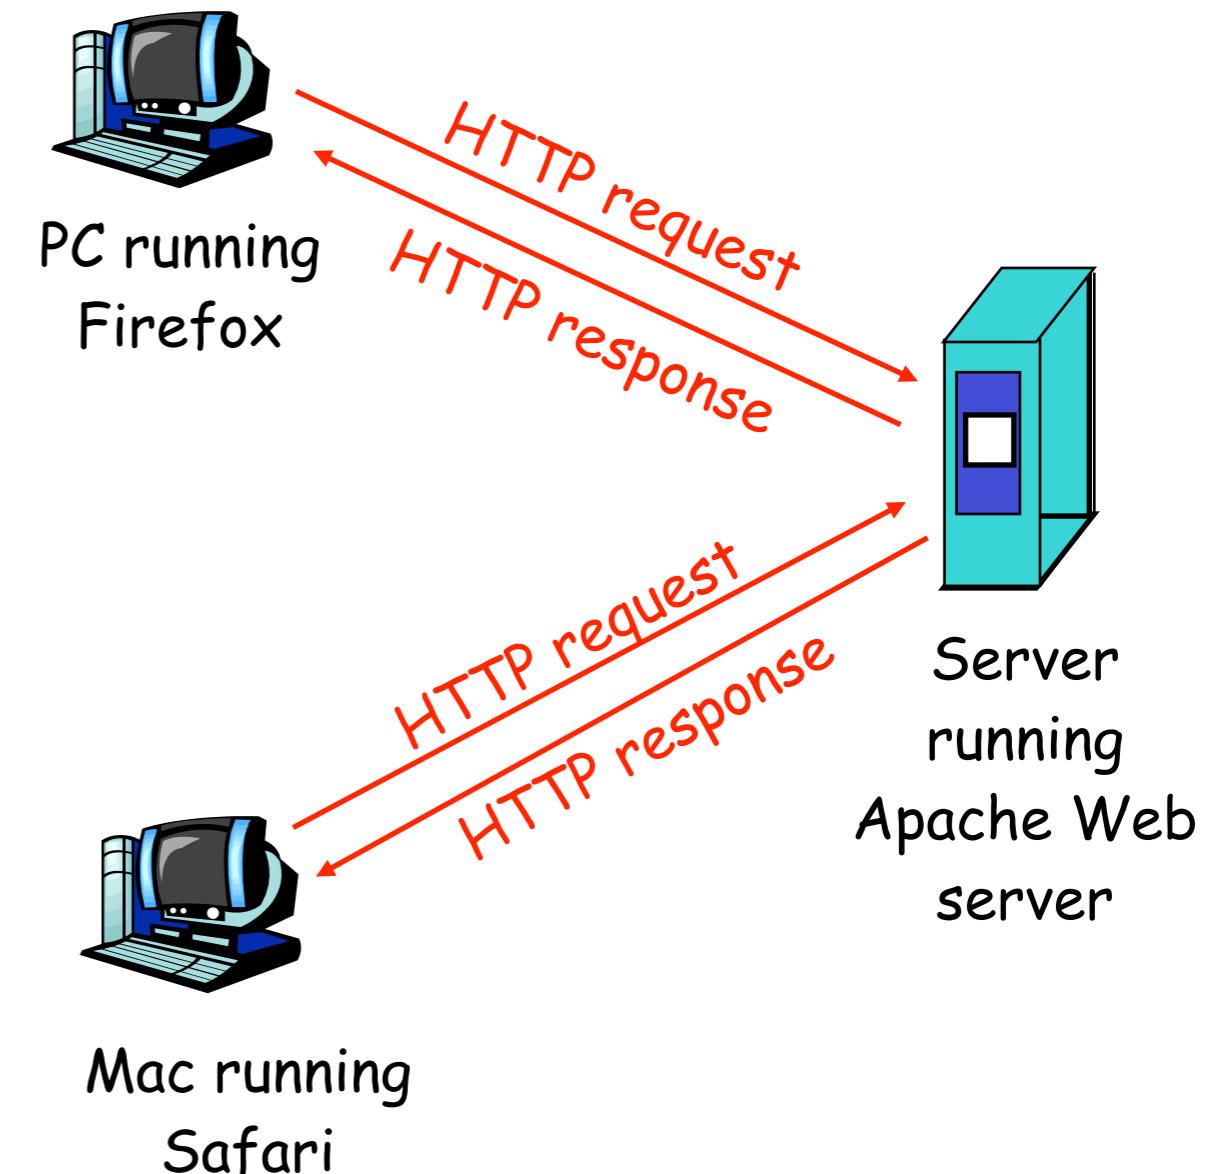
- Web's application layer protocol (RFCs 1945, 2616)
- client/server model
  - **client:** browser that requests, receives, “displays” Web objects
  - **server:** Web server sends objects in response to requests



# HTTP overview

## HTTP: hypertext transfer protocol

- Web's application layer protocol (RFCs 1945, 2616)
- client/server model
  - **client:** browser that requests, receives, “displays” Web objects
  - **server:** Web server sends objects in response to requests



Question: What programming language used to code client/server?

# HTTP overview (continued)

## Uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed
  - Well, maybe (see next slide)

## HTTP is “stateless”

- server maintains no information about past client requests

aside

## Protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP connections

## Nonpersistent HTTP

- At most one object is sent over a TCP connection.
- HTTP/1.0 uses nonpersistent HTTP

## Persistent HTTP

- Multiple objects can be sent over single TCP connection between client and server.
- HTTP/1.1+ use persistent connections in default mode



# Nonpersistent HTTP

Suppose user enters URL

www.someSchool.edu/someDepartment/home.index

(contains text,  
references to 10  
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at www.someSchool.edu on port 80

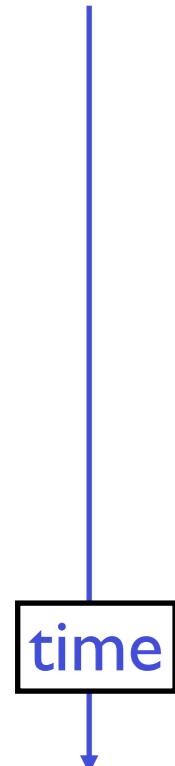
1b. HTTP server at host www.someSchool.edu waiting for TCP connection at port 80. “accepts” connection, notifying client

2. HTTP client sends HTTP **request message** (containing URL) into TCP connection socket. Message indicates that client wants object someDepartment/home.index

3. HTTP server receives request message, forms **response message** containing requested object, and sends message into its socket

time

# Nonpersistent HTTP (cont.)



5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects
6. Steps 1-5 repeated for each of 10 jpeg objects

4. HTTP server closes TCP connection.



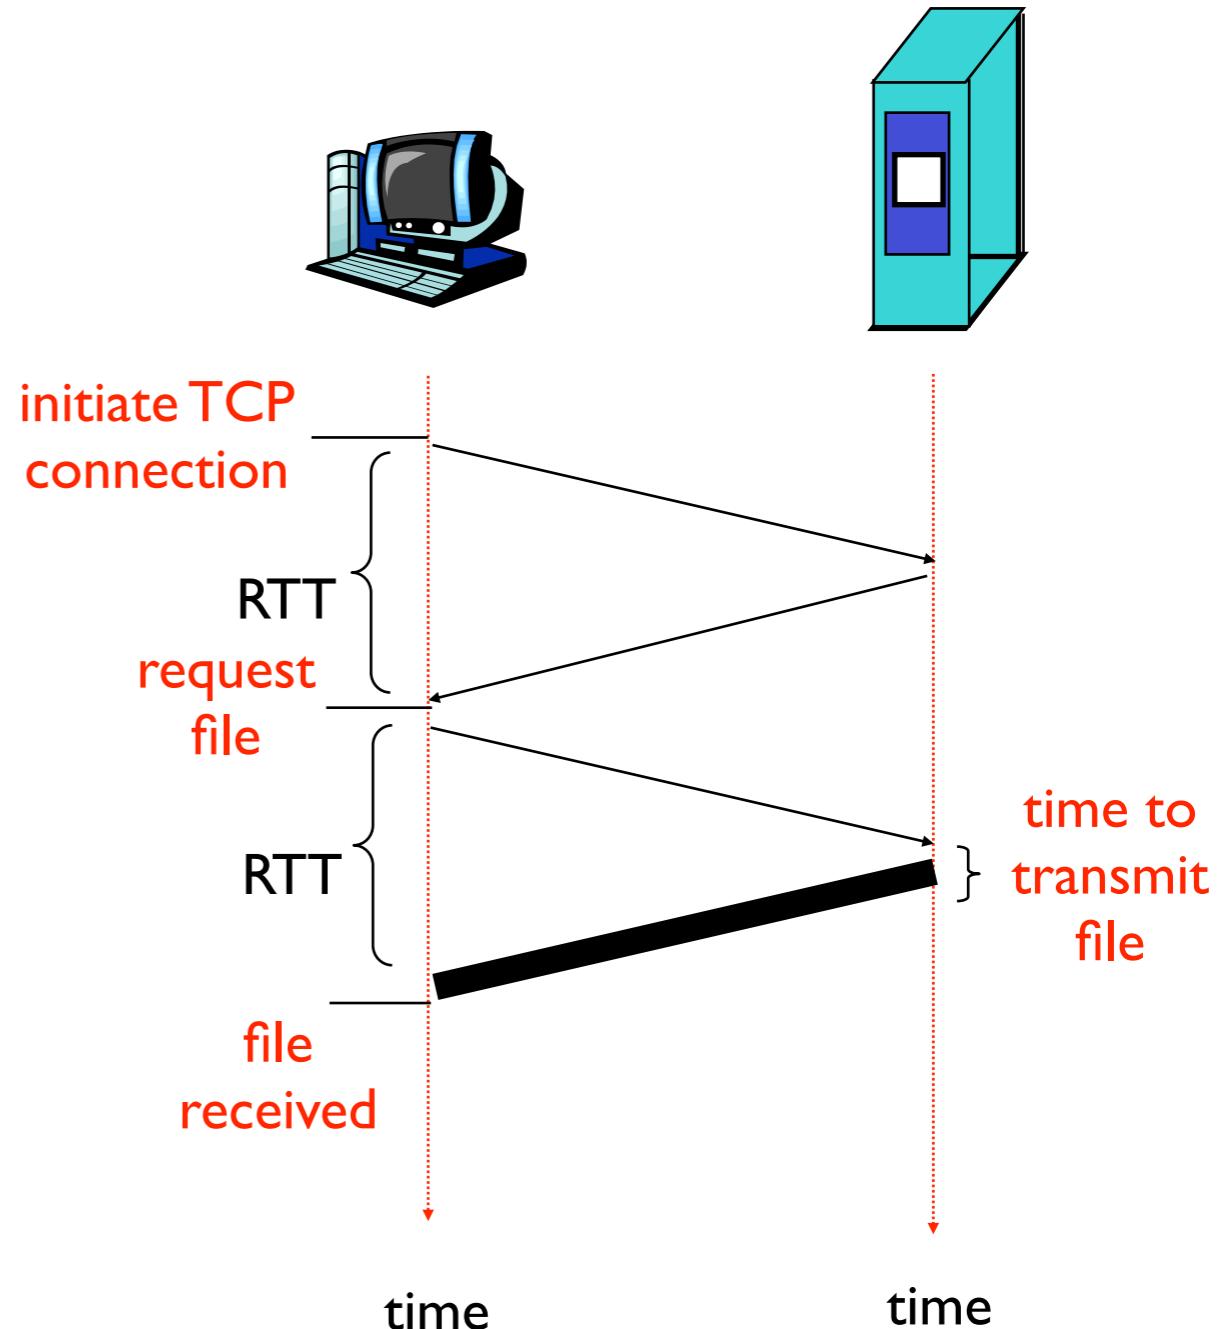
# Non-Persistent HTTP: Response time

**Definition of RTT:** time to send a small packet to travel from client to server and back.

## Response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

$$\text{total} = 2\text{RTT} + \text{transmit time}$$



# Persistent HTTP

## Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection

## Persistent without pipelining:

- client issues new request only when previous response has been received
- one RTT for each referenced object

## Persistent with pipelining:

- default in HTTP/1.1
- client sends requests as soon as it encounters a referenced object
- as little as one RTT total for all the referenced objects

# Persistent HTTP

## Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection

## Persistent without pipelining:

- client issues new request only when previous response has been received
- one RTT for each referenced object

## Persistent with multiplexing:

- default in HTTP/2
- client interleaves requests (why is this better?)
- allows for true parallelism

# Persistent vs Non-Persistent HTTP

- Given advantages of non-persistent HTTP, why use non-persistent?
  - ▶ Possibly scalability issue: If users have connections open but aren't using them, others may not be able to connect
  - ▶ Simplicity: non-persistent is easier to deal with than concurrency issues with pipelined persistent HTTP



Regardless, as non-persistent is not the default since V1, advantages don't outweigh disadvantages

# HTTP request message

- two types of HTTP messages: **request, response**
- **HTTP request message:**

- ASCII (human-readable format)

request line  
(GET, POST,  
HEAD commands)

header lines

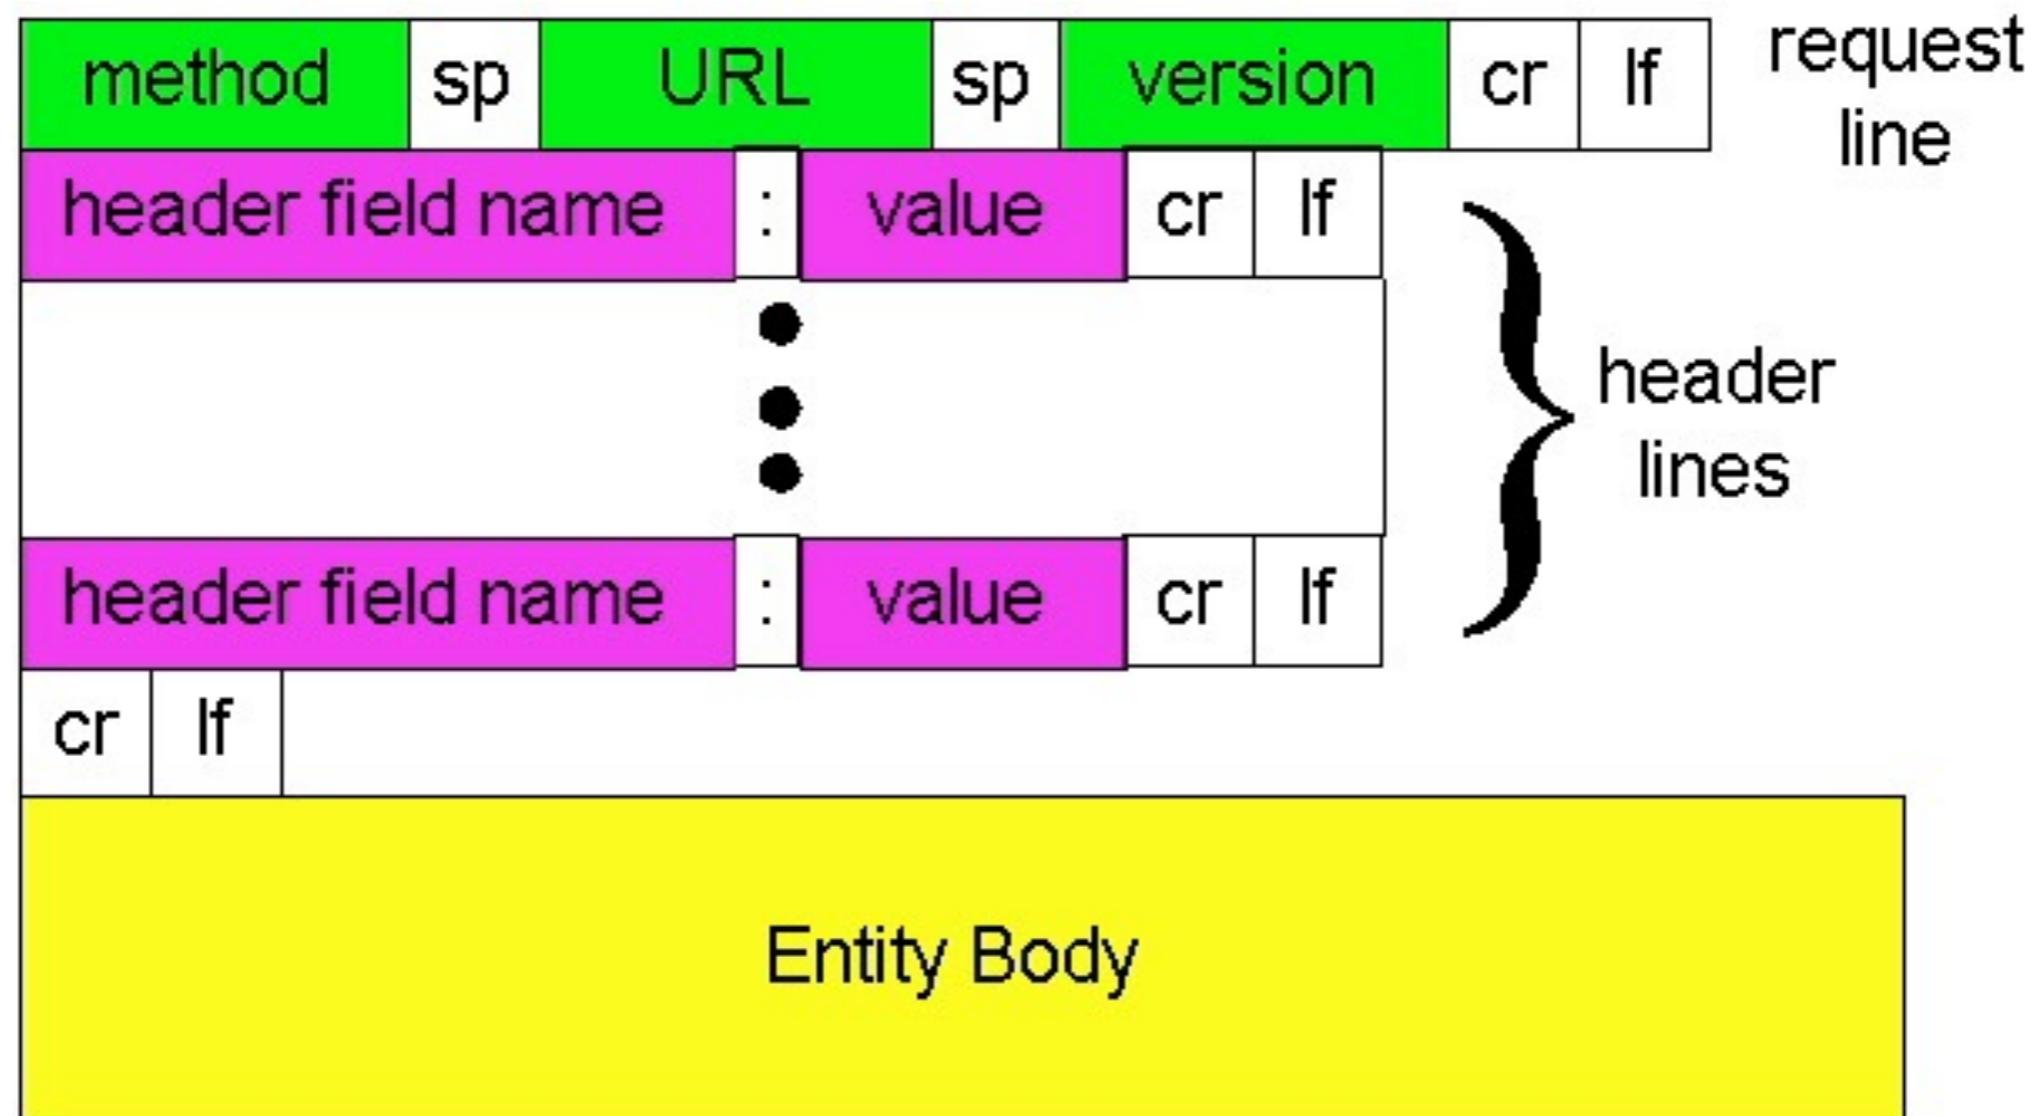
Carriage return,  
line feed  
indicates end  
of message

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr


```

(extra carriage return, line feed)

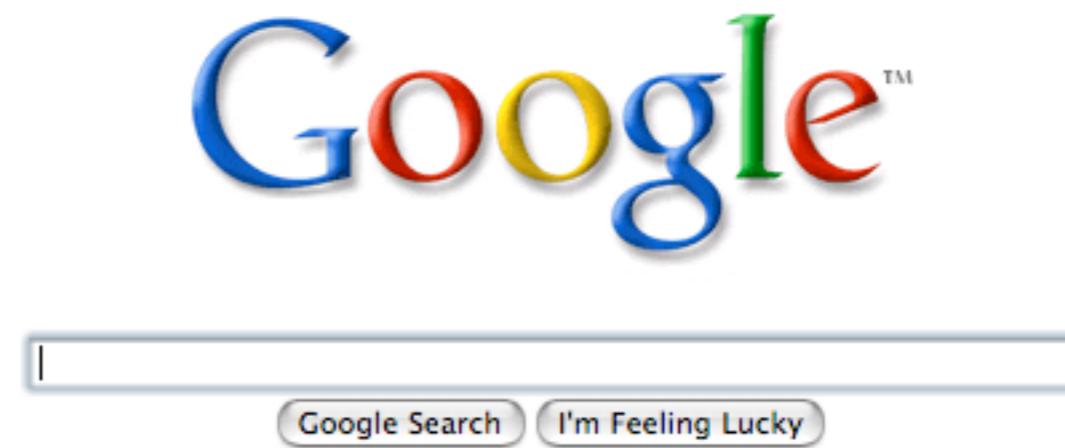
# HTTP request message: general format



# Uploading form input

## Post method:

- Web page often includes form input
- Input is uploaded to server in entity body



## URL method:

- Uses GET method
- Input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

# Method types

## HTTP/1.0

- GET
- POST
- HEAD
  - asks server to leave requested object out of response

## HTTP/1.1

- GET, POST, HEAD
- PUT
  - uploads file in entity body to path specified in URL field
- DELETE
  - deletes file specified in the URL field

# Method types

PUT: The Request-URI in a PUT request identifies the target resource itself. The client specifies the exact location (URI) where the enclosed entity should be stored or updated. If a resource already exists at that URI, it is replaced; otherwise, a new one is created.

POST: The Request-URI in a POST request identifies the resource that will handle the enclosed entity. This often means the URI points to a collection, and the server is responsible for creating a new resource within that collection, assigning it a new URI. The client does not specify the exact URI of the newly created resource.

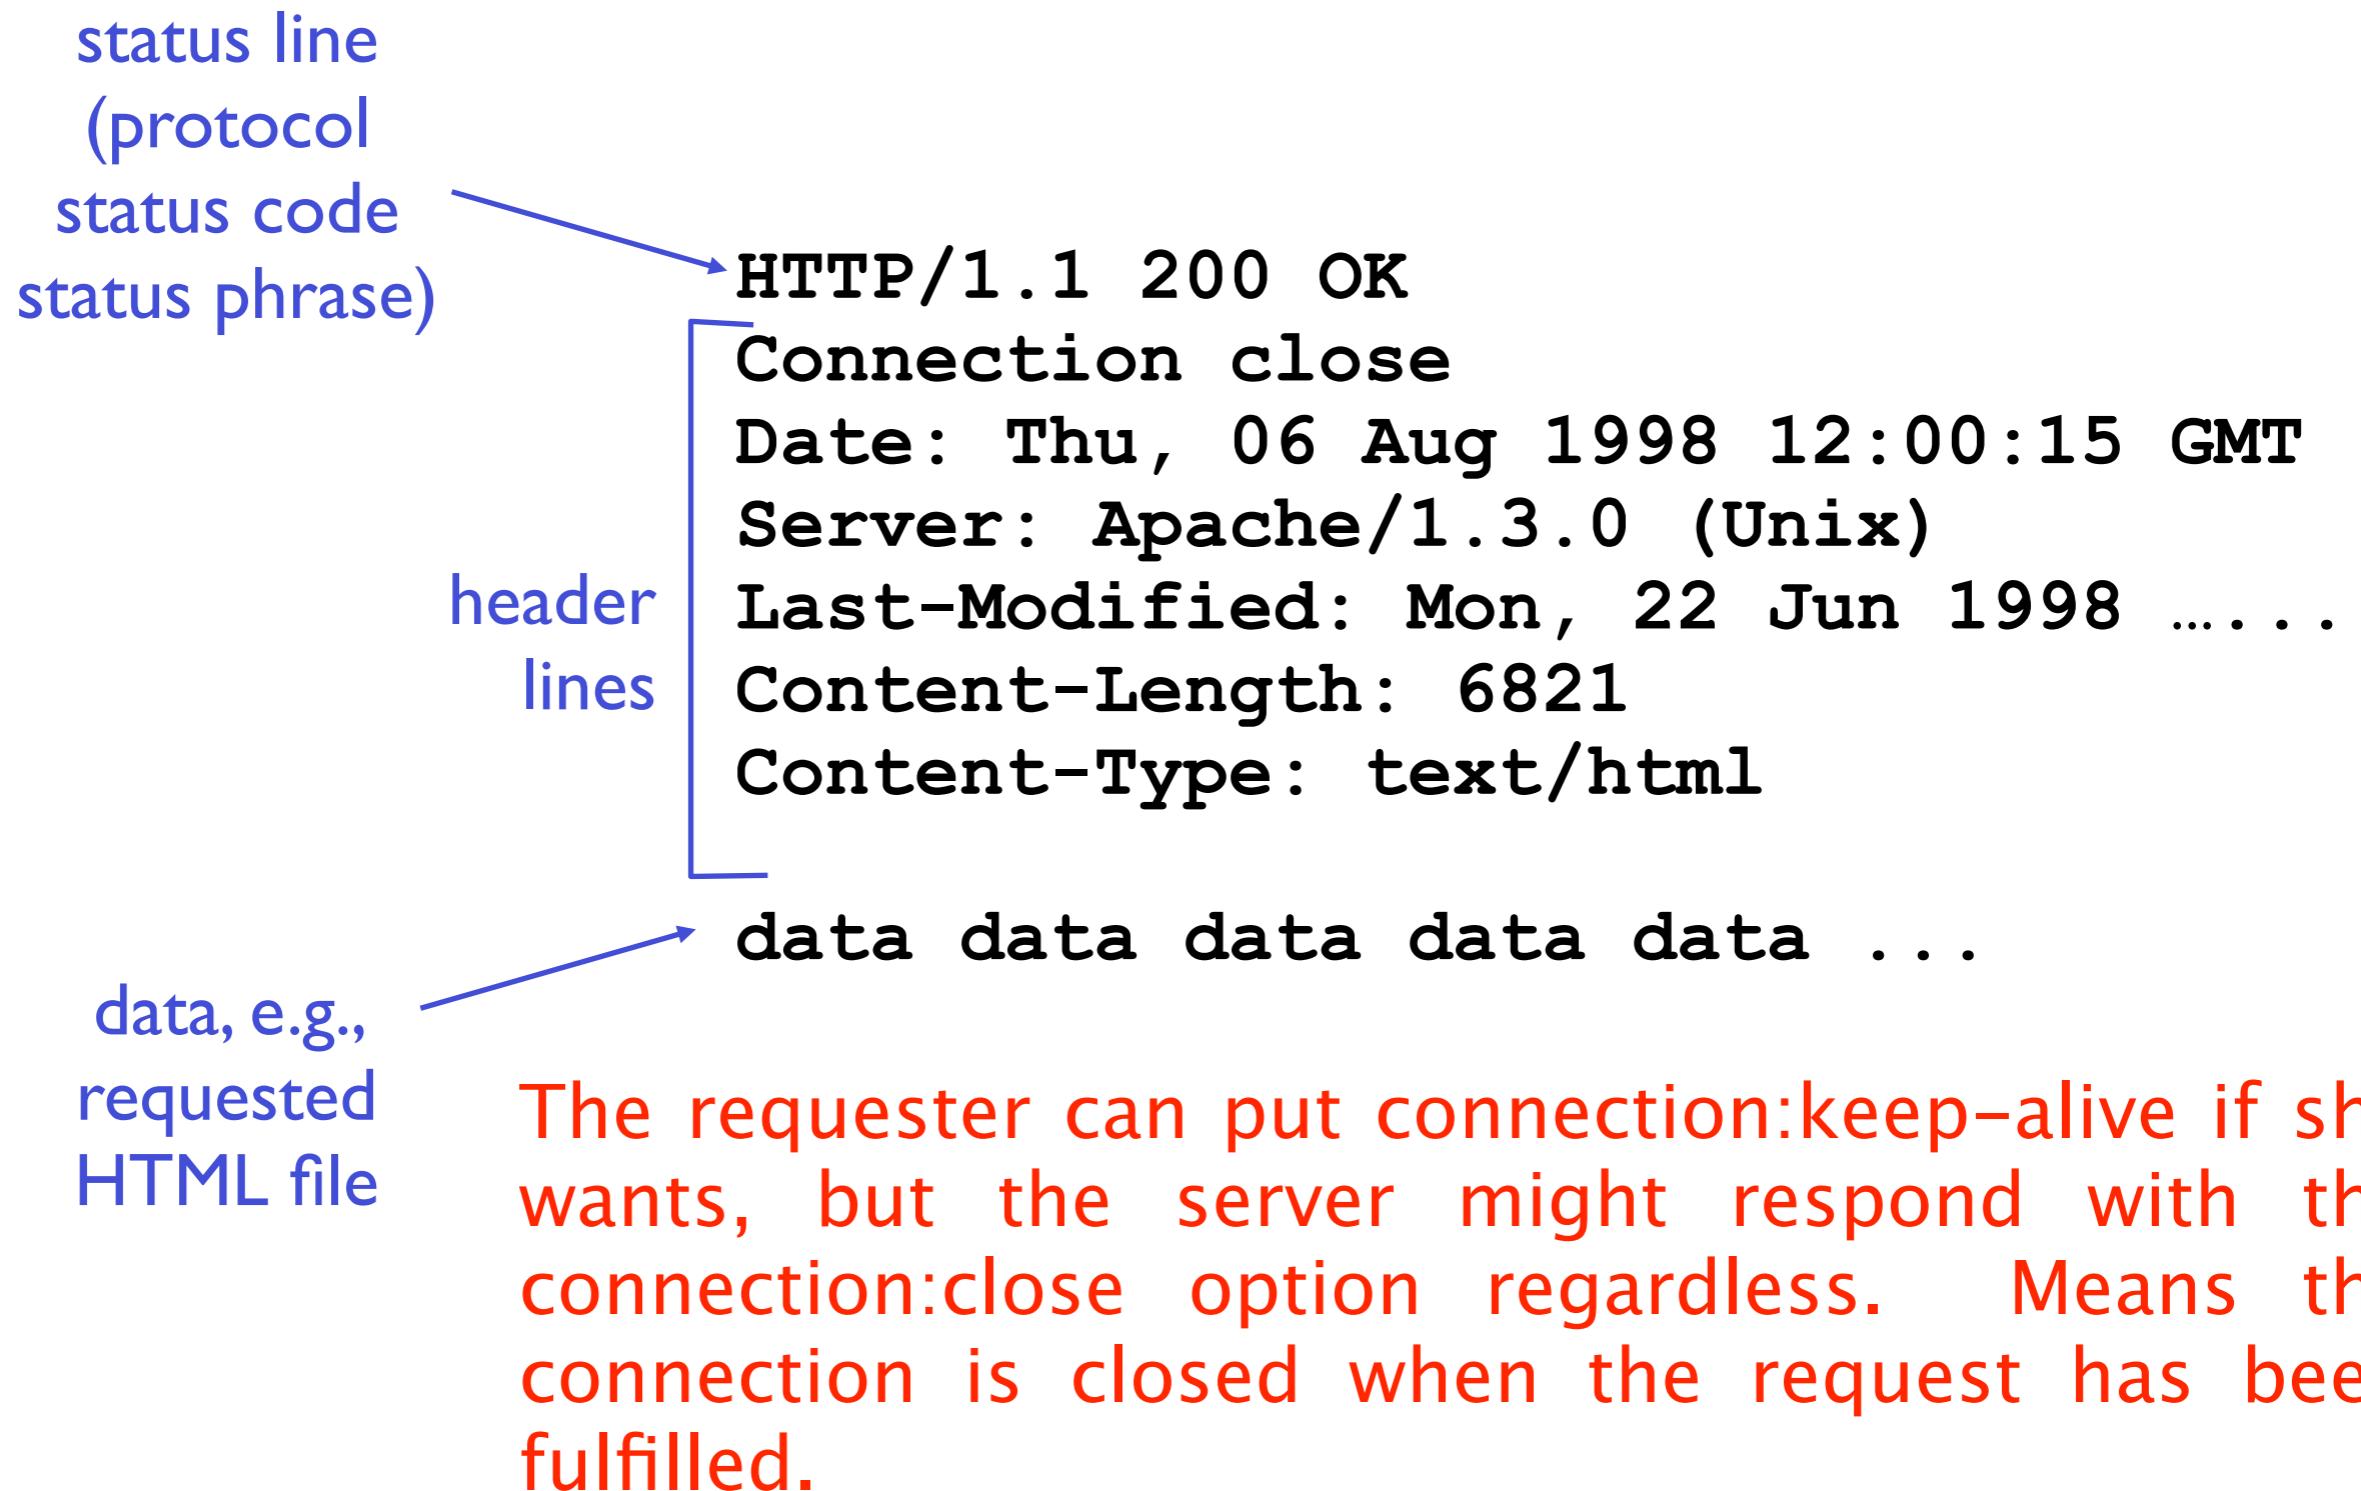
# PUT vs POST

- From the RFC: The fundamental difference between the POST and PUT requests is reflected in the different meaning of the Request-URI. The URI in a POST request **identifies the resource that will handle the enclosed entity**. That resource might be a data-accepting process, a gateway to some other protocol, or a separate entity that accepts annotations. In contrast, the URI in a PUT request **identifies the entity enclosed with the request** -- the user agent knows what URI is intended and the server **MUST NOT** attempt to apply the request to some other resource. If the server desires that the request be applied to a different URI, it **MUST** send a 301 (Moved Permanently) response; the user agent **MAY** then make its own decision regarding whether or not to redirect the request.

# PUT vs POST

Bottom line: POST gives data to an entity at a specific URI and expects that entity to handle it. PUT puts a file or resource at a specific URI. Also differences about idempotency.

# HTTP response message



# HTTP response status codes

In the first line in the server-to-client response message.

A few sample codes:

## 200 OK

- request succeeded, requested object later in this message

## 301 Moved Permanently

- requested object moved, new location specified later in this message (Location:)

## 400 Bad Request

- request message not understood by server

## 404 Not Found

- requested document not found on this server

## 505 HTTP Version Not Supported

# Trying out HTTP (client side) for yourself

## I. Telnet to your favorite Web server:

```
telnet www.richmond.edu 80
```

Opens TCP connection to port 80  
(default HTTP server port) at  
www.richmond.edu.

Anything typed in sent  
to port 80 at www.richmond.edu

## 2. Type in a GET HTTP request:

```
GET /~dszajda/classes/cs332/  
Fall_2019/index.html HTTP/1.1  
Host: www.richmond.edu
```

By typing this in (hit carriage  
return twice), you send  
this minimal (but complete)  
GET request to HTTP server

## 3. Look at response message sent by HTTP server!

# User-server state: cookies

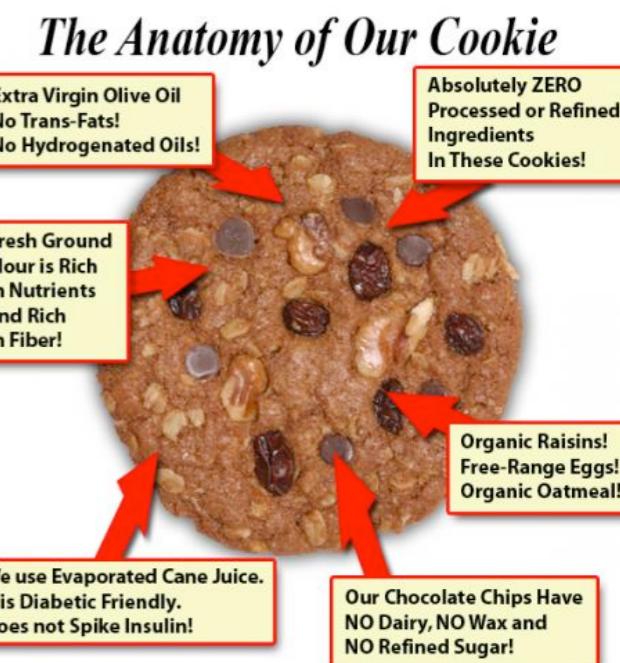
Many major Web sites use cookies

## Example:

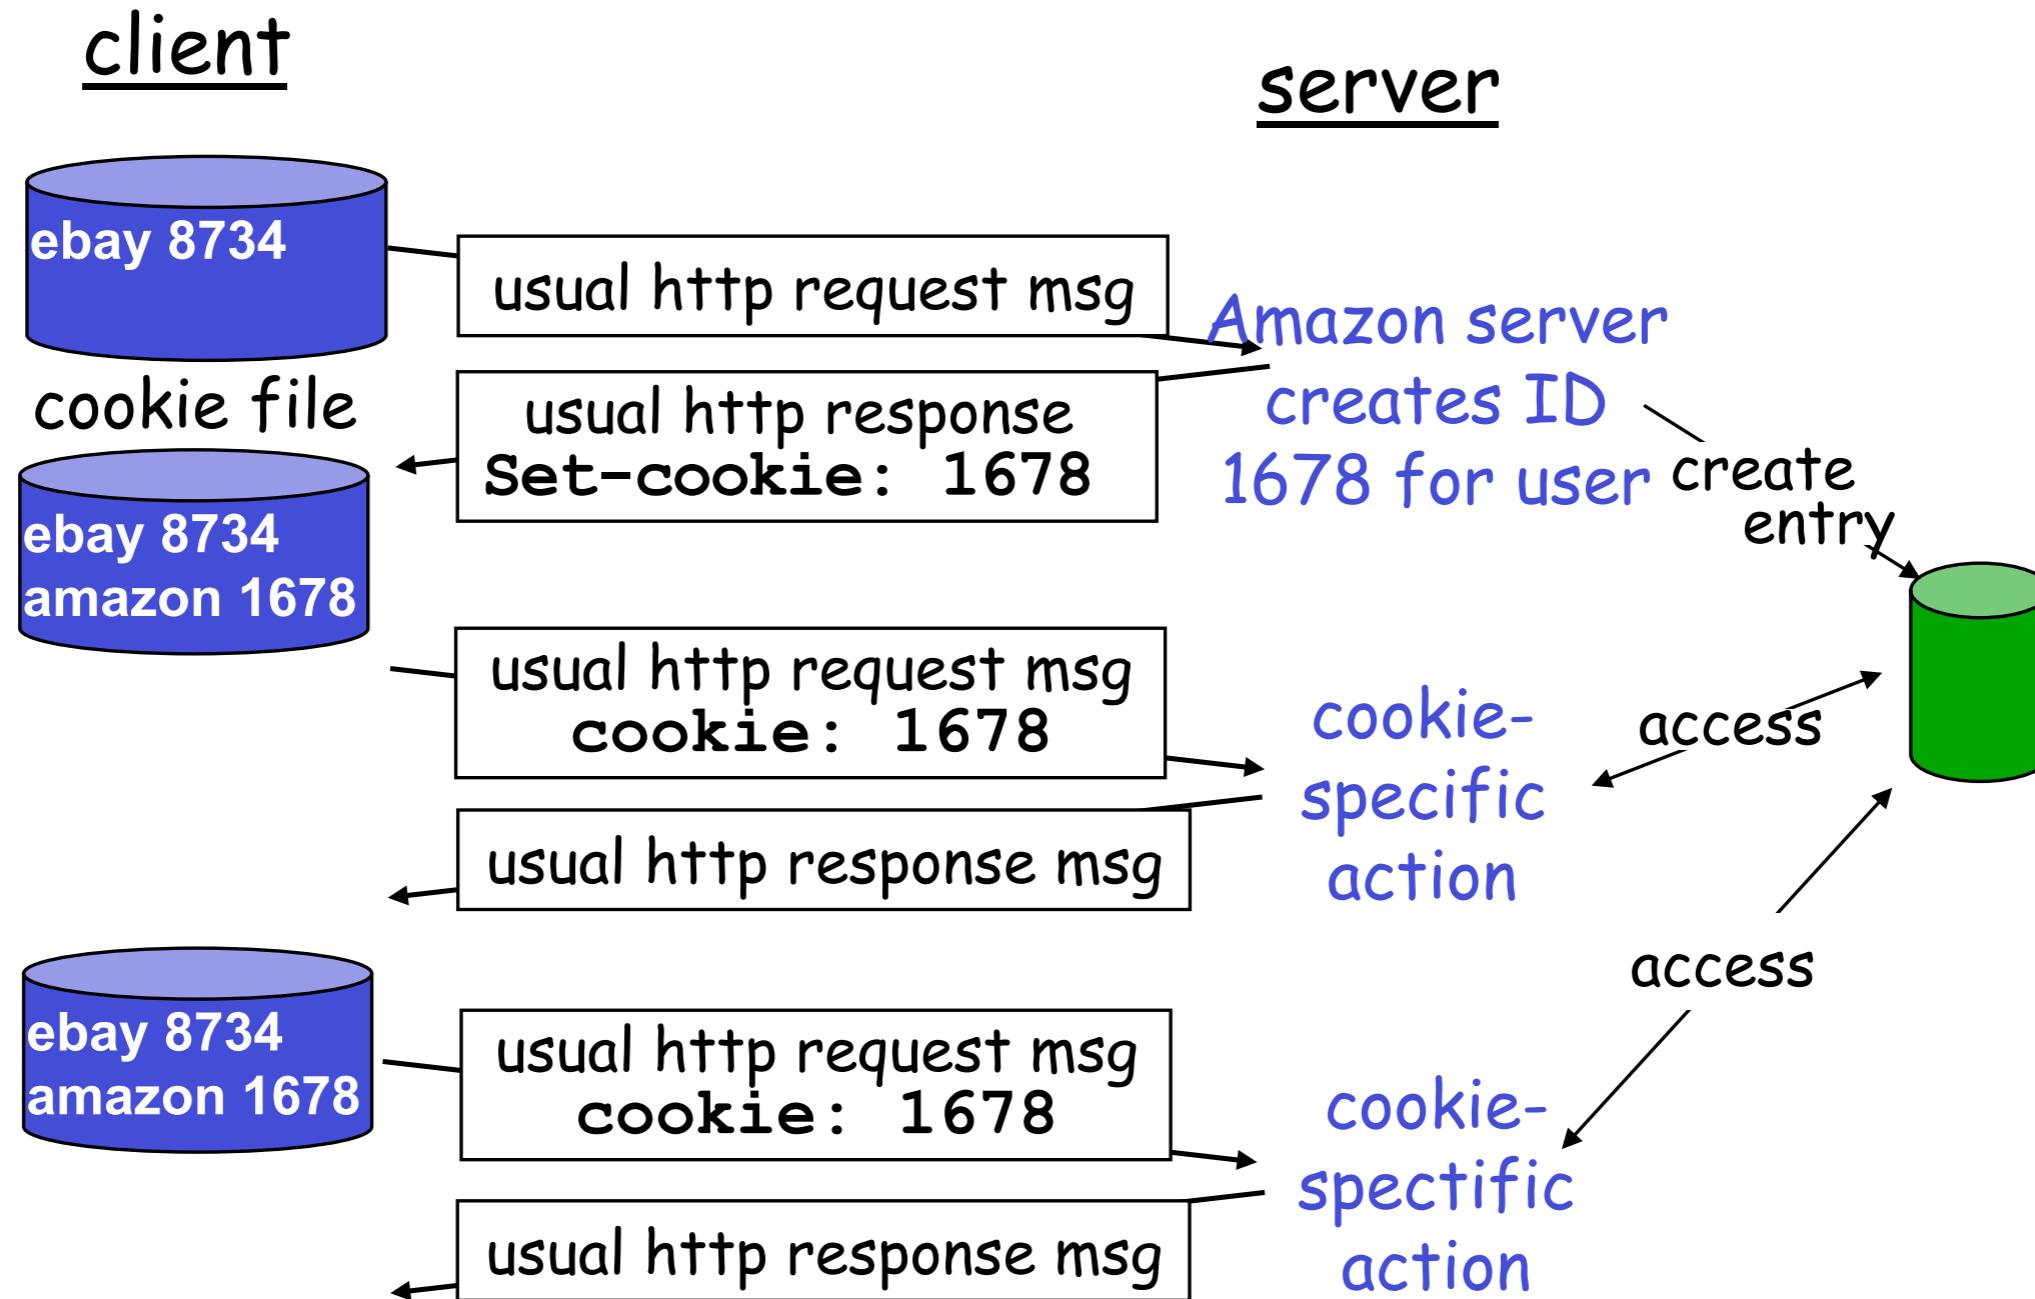
### Four components:

- 1) cookie header line of HTTP response message
- 2) cookie header line in HTTP request message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

- ▶ Susan accesses Internet always from same PC
- ▶ She visits a specific e-commerce site for first time
- ▶ When initial HTTP requests arrives at site, site creates a unique ID and creates an entry in backend database for ID



# Cookies: keeping “state” (cont.)



# Cookies (continued)

## What cookies can bring:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

aside

## Cookies and privacy:

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

## How to keep “state”:

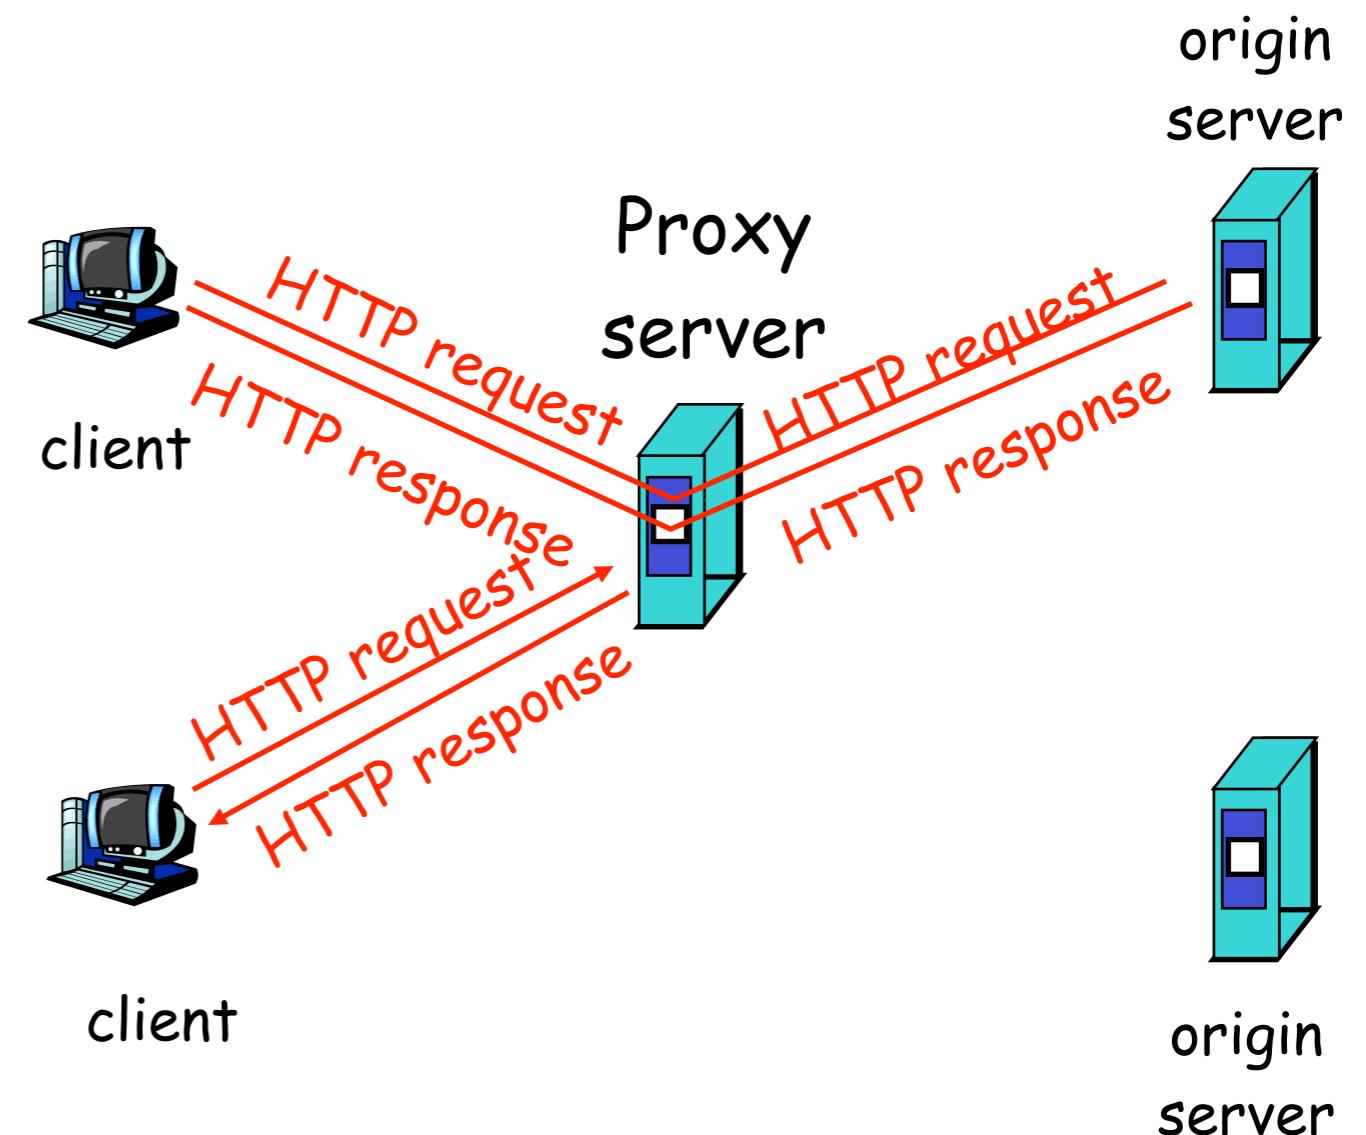
- Protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state



# Web caches (proxy server)

**Goal:** satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
  - ▶ object in cache: cache returns object
  - ▶ else cache requests object from origin server, then returns object to client



# More about Web caching

- Cache acts as both client and server
- Typically cache is installed by ISP (university, company, residential ISP)



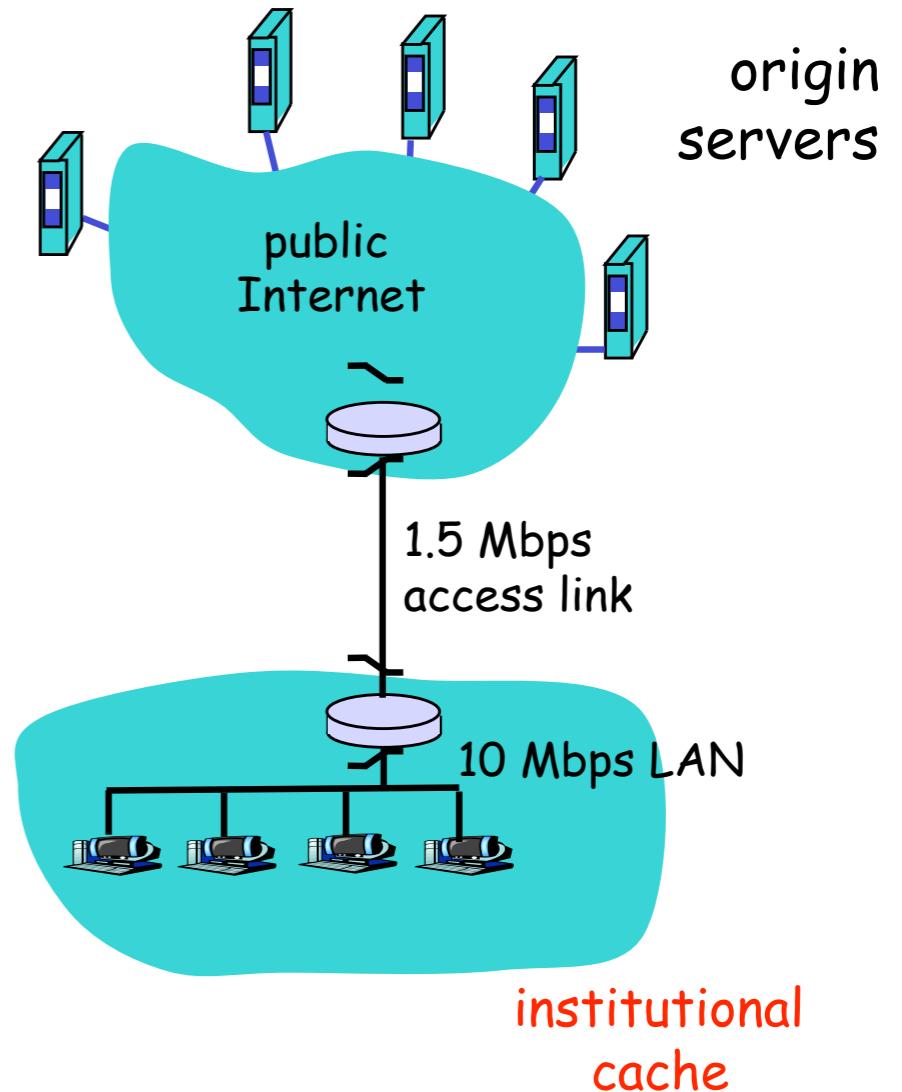
## Why Web caching?

- Reduce response time for client request.
- Reduce traffic on an institution's access link.
- Internet dense with caches: enables “poor” content providers to effectively deliver content (but so does P2P file sharing)

# Question: What to do?

## Assumptions

- average object size = 100,000 bits
- avg. request rate from institution's browsers to origin servers = 15/sec
- delay from institutional router to any origin server and back to router = 2 sec



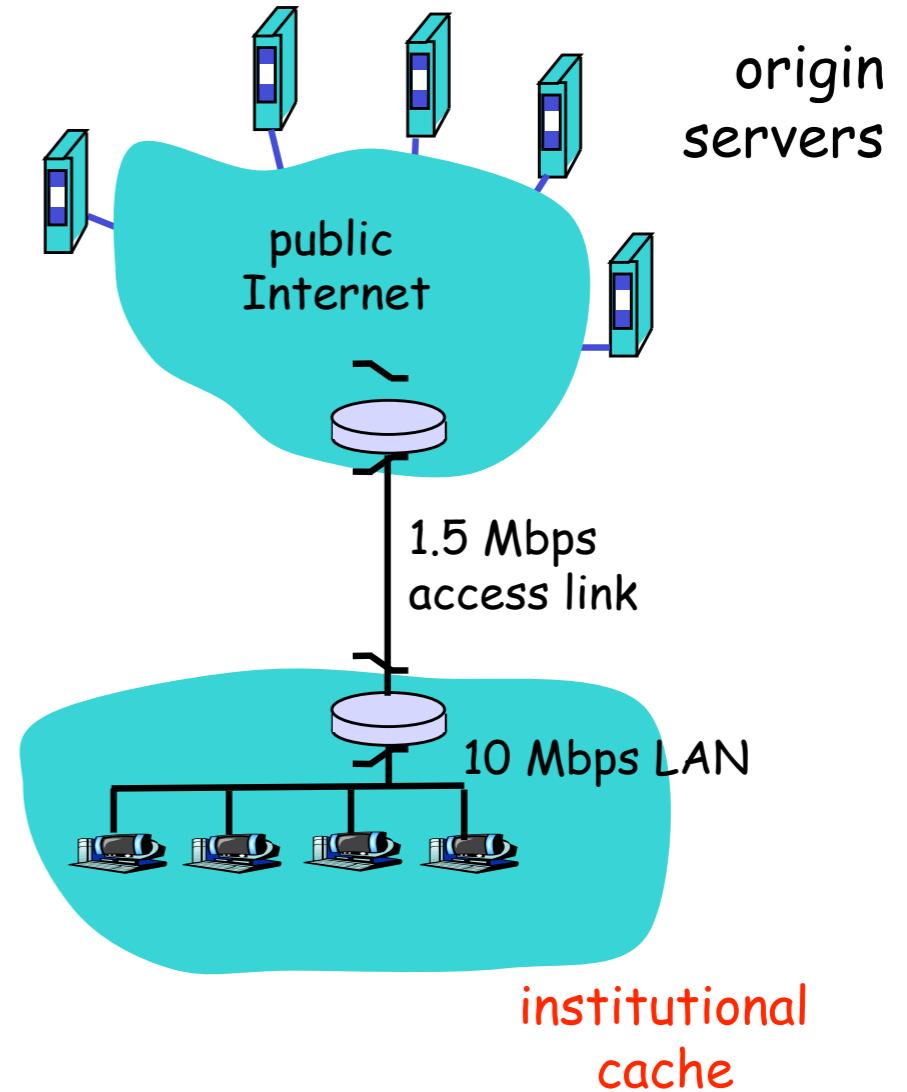
# Caching example

## Assumptions

- average object size = 100,000 bits
- avg. request rate from institution's browsers to origin servers = 15/sec
- delay from institutional router to any origin server and back to router = 2 sec

## Consequences

- utilization on LAN = 15%
- utilization on access link = 100%
- total delay = Internet delay + access delay + LAN delay  
 $= 2 \text{ sec} + \text{minutes} + \text{milliseconds}$



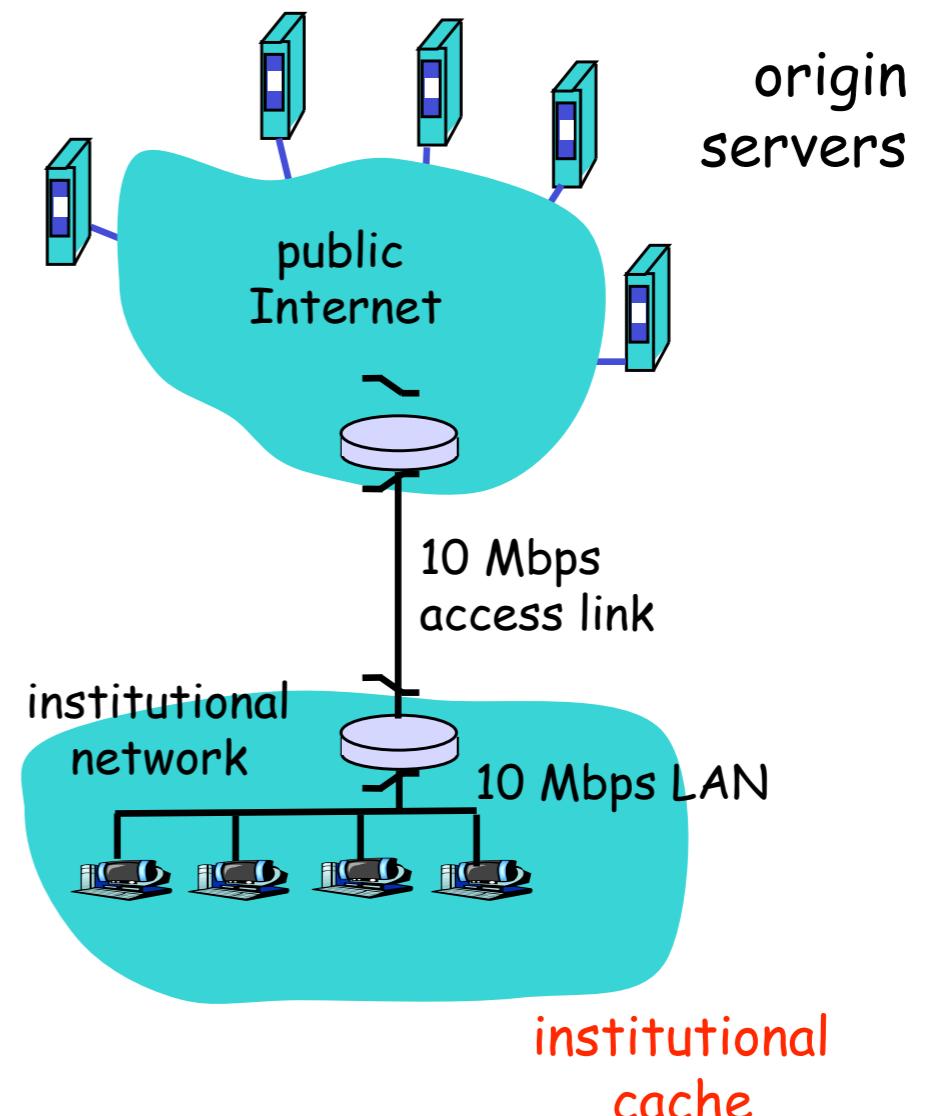
# Caching example (cont)

## Possible solution

- increase bandwidth of access link to, say, 10 Mbps

## Consequences

- utilization on LAN = 15%
- utilization on access link = 15%
- Total delay = Internet delay + access delay + LAN delay  
 $= 2 \text{ sec} + \text{msecs} + \text{msecs}$
- often a costly upgrade



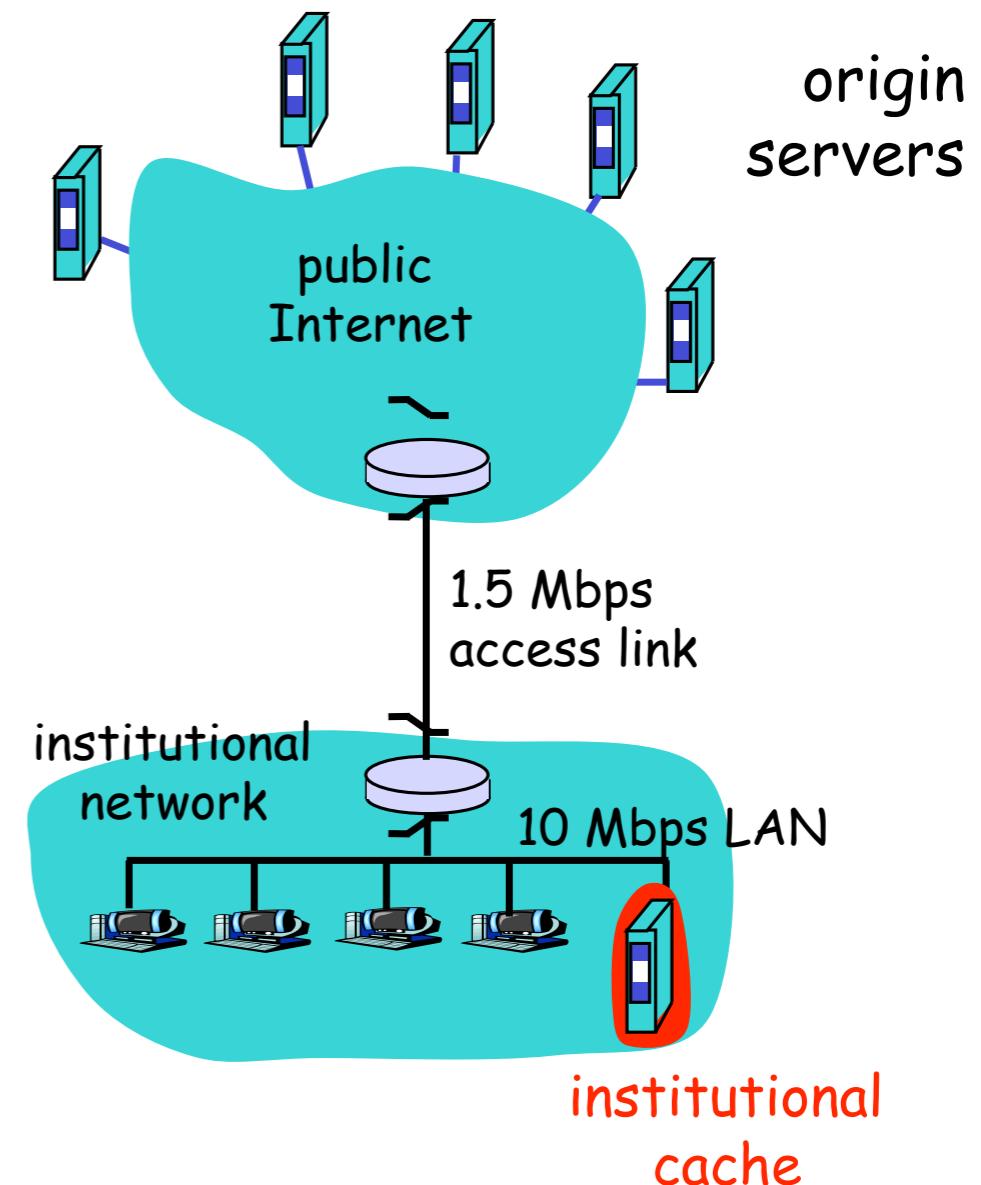
# Caching example (cont)

## Install cache

- suppose hit rate is .4

## Consequence

- 40% requests will be satisfied almost immediately
- 60% requests satisfied by origin server
- utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- total avg delay = Internet delay + access delay + LAN delay =  $0.6*(2.01) \text{ secs} + .4*\text{milliseconds} < 1.4 \text{ secs}$



# Conditional GET

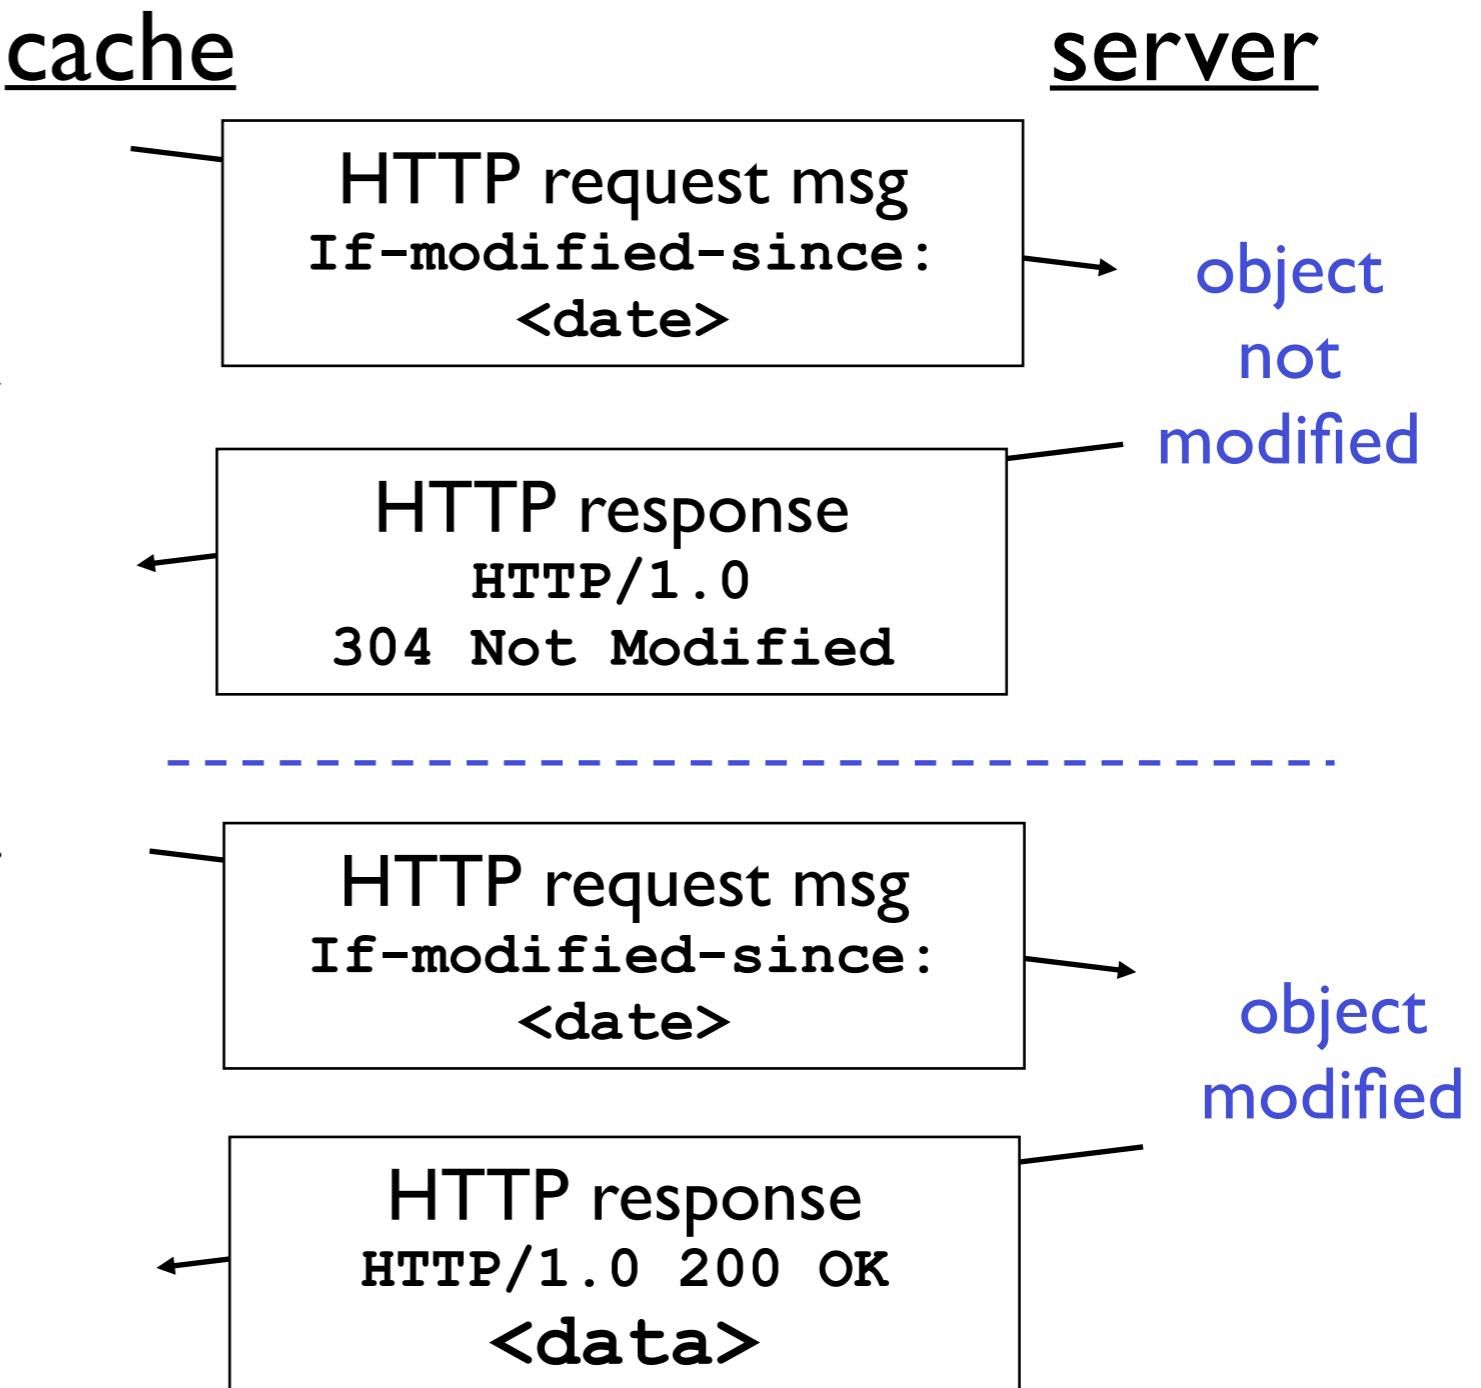
- **Goal:** don't send object if cache has up-to-date cached version

- cache: specify date of cached copy in HTTP request

**If-modified-since: <date>**

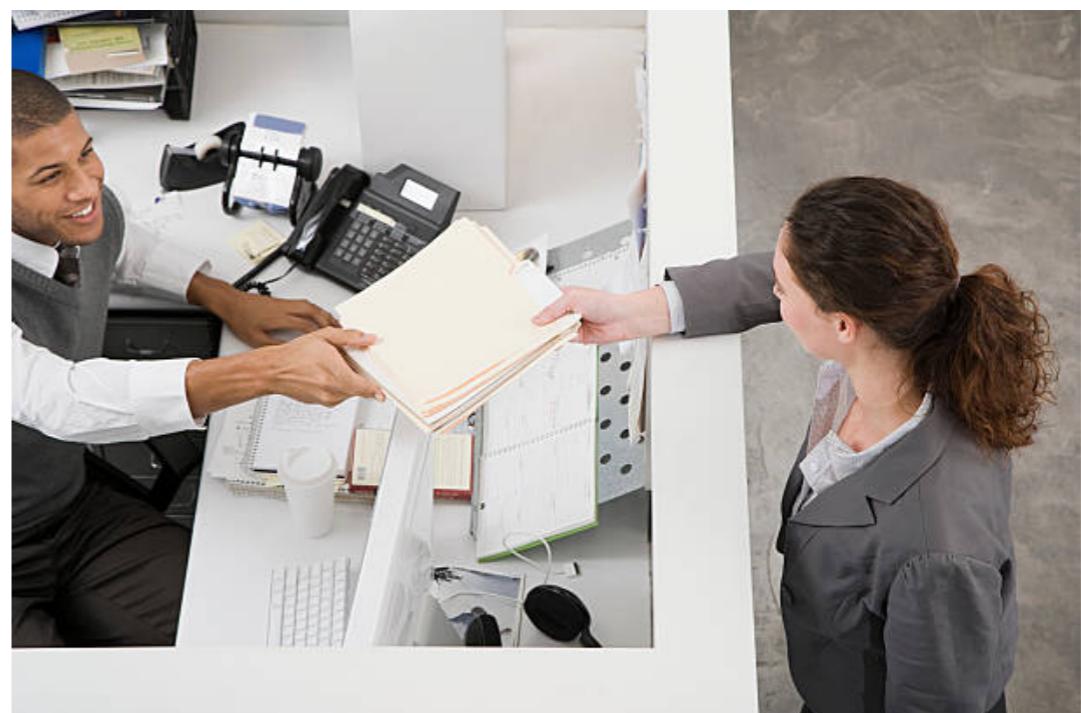
- server: response contains no object if cached copy is up-to-date:

**HTTP/1.0 304 Not Modified**

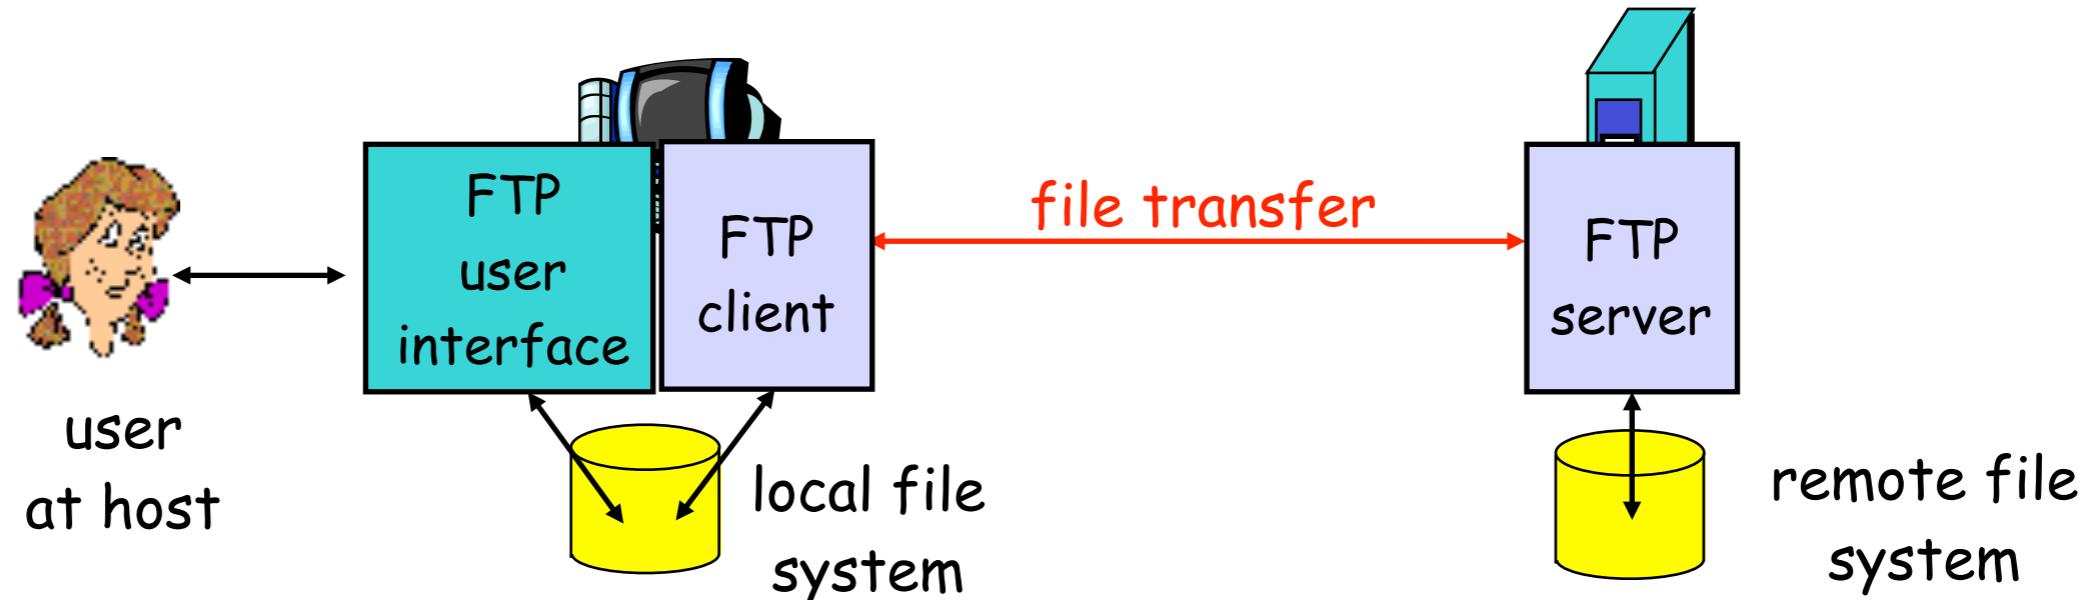


# Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
- 2.5 DNS
- 2.6 P2P file sharing
- 2.7-2.8 Sockets
- 2.9 Building a webserver

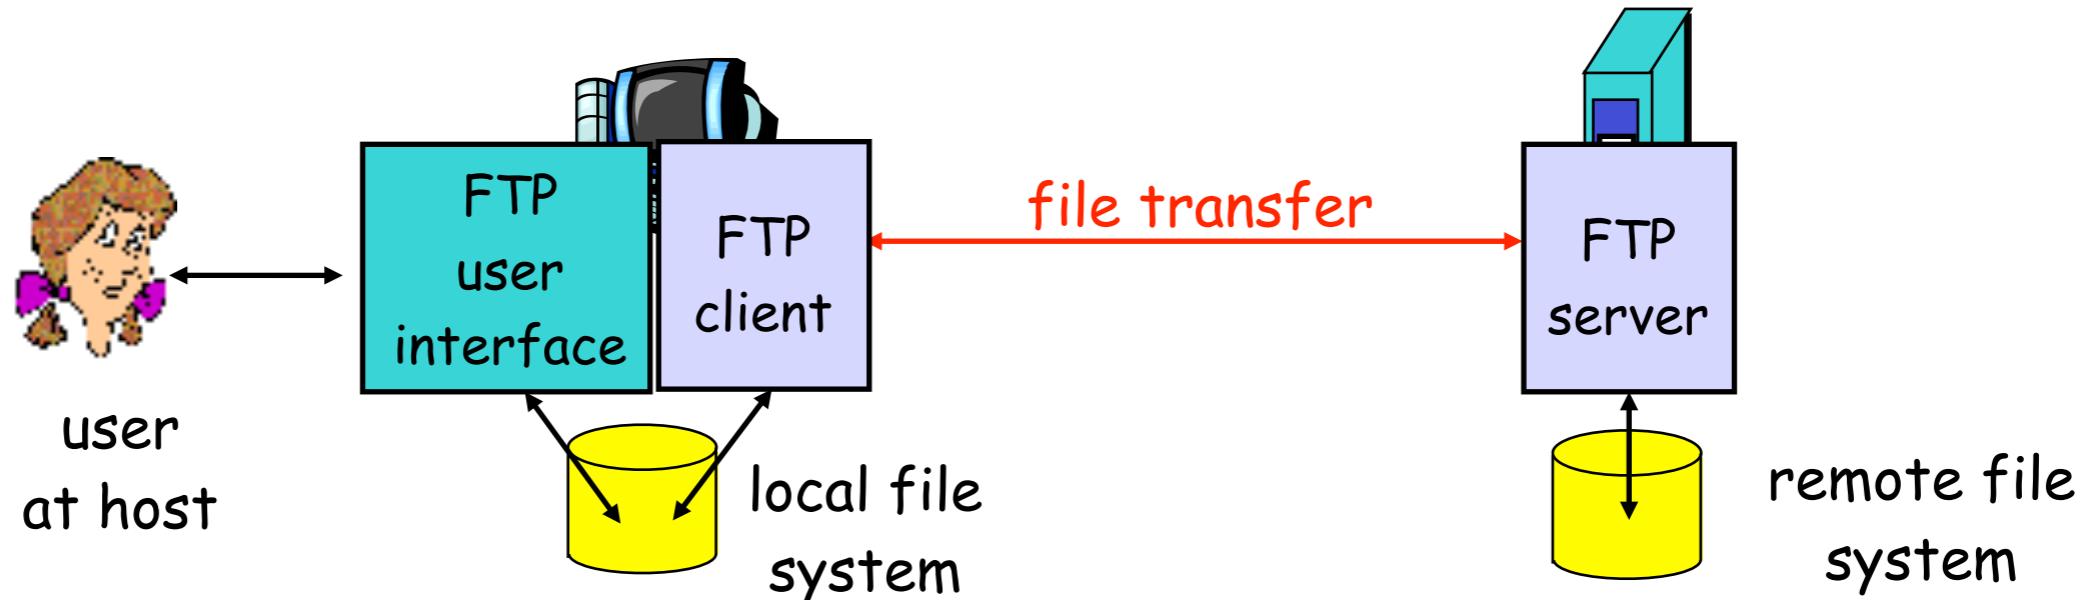


# FTP: the file transfer protocol



- transfer file to/from remote host
- You design it
  - ▶ What kind of messages do you need
  - ▶ What kind of responses should a given message generate?

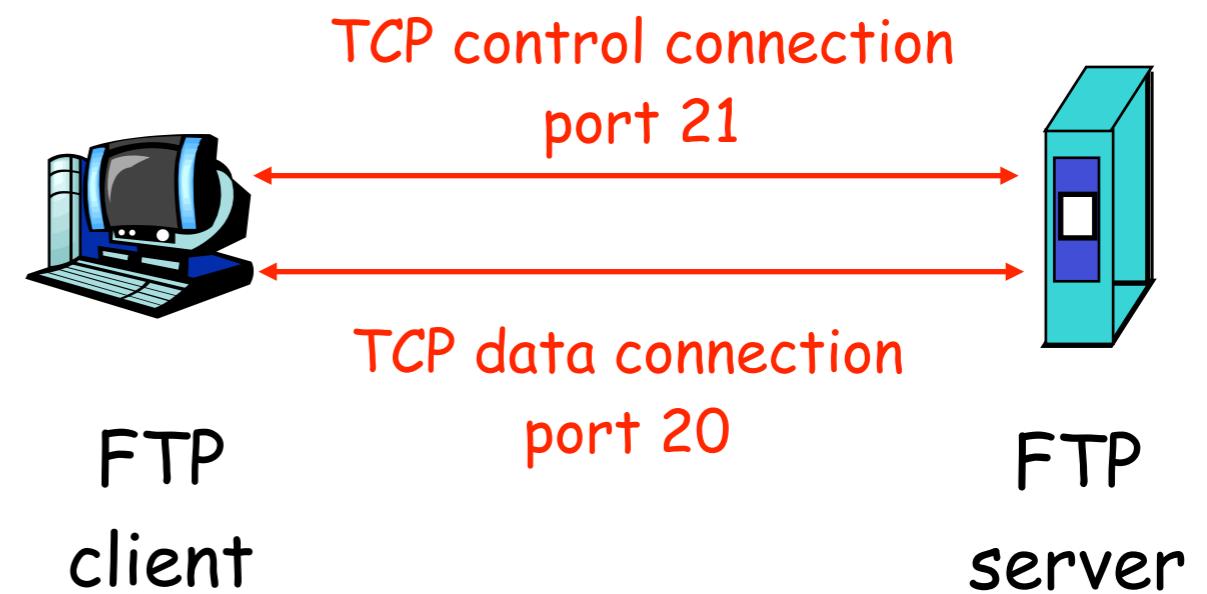
# FTP: the file transfer protocol



- transfer file to/from remote host
- client/server model
  - ▶ **client:** side that initiates transfer (either to/from remote)
  - ▶ **server:** remote host
- ftp: RFC 959
- ftp server: port 21

# FTP: separate control, data connections

- FTP client contacts FTP server at port 21, specifying TCP as transport protocol
- Client obtains authorization over control connection
- Client browses remote directory by sending commands over control connection.
- When server receives file transfer command, server opens 2<sup>nd</sup> TCP connection (for file) to client
- After transferring one file, server closes data connection.



- Server opens another TCP data connection to transfer another file.
- Control connection: “out of band”
- FTP server maintains “state”: current directory, earlier authentication

# FTP commands, responses

## Sample commands:

- sent as ASCII text over control channel
- **USER *username***
- **PASS *password***
- **LIST** return list of file in current directory
- **RETR *filename*** retrieves (gets) file
- **STOR *filename*** stores (puts) file onto remote host

## Sample return codes

- status code and phrase (as in HTTP)
- **331 Username OK, password required**
- **125 data connection already open; transfer starting**
- **425 Can't open data connection**
- **452 Error writing file**