

CMSC 332

Computer Networks

Transport Layer

Professor Szajda

Chapter 3: Transport Layer

Our goals:

- understand principles behind transport layer services:
 - ▶ multiplexing/demultiplexing
 - ▶ reliable data transfer
 - ▶ flow control
 - ▶ congestion control
- learn about transport layer protocols in the Internet:
 - ▶ UDP: connectionless transport
 - ▶ TCP: connection-oriented transport
 - ▶ TCP congestion control




Chapter 3 Outline

- **3.1 Transport-layer services**
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - ▶ segment structure
 - ▶ reliable data transfer
 - ▶ flow control
 - ▶ connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
 - ▶ send side: breaks app messages into *segments*, passes to network layer
 - ▶ rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
 - ▶ Internet: TCP and UDP



Transport vs. Network layer

- **network layer:** logical communication between **hosts**
- **transport layer:** logical communication between **processes**
 - ▶ relies on, enhances, network layer services



Household analogy:

- 12 kids sending letters to 12 kids
- processes = kids
 - app messages = letters in envelopes
 - hosts = houses
 - transport protocol = Ann and Bill
 - network-layer protocol = postal service


Layers of Networks?

- You can view each layer that we have discussed thus far as an abstract network:
 - ▶ Application Layer Networks: P2P, Social Networks, etc
 - ▶ Transport Layer Networks: Communicating processes
 - ▶ Network Layer Networks: Networks of Hosts
 - ▶ Link Layer Networks: One-Hop Networks
 - ▶ Physical Layer Networks: Wires



Internet transport-layer protocols

- reliable, in-order delivery,
no duplicates (TCP)
 - ▶ congestion control
 - ▶ flow control
 - ▶ connection setup
- unreliable, unordered
delivery: UDP
 - ▶ no-frills extension of “best-effort” IP
- services not available:
 - ▶ delay guarantees
 - ▶ bandwidth guarantees



Chapter 3 Outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - ▶ segment structure
 - ▶ reliable data transfer
 - ▶ flow control
 - ▶ connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

Multiplexing/demultiplexing

Demultiplexing at rcv host:




delivering received segments
to correct socket

Multiplexing at send host:

gathering data from multiple
sockets, enveloping data with
header (later used for
demultiplexing)

= socket

= process




host 1

host 2

host 3

How demultiplexing works

- host receives IP datagrams
 - ▶ each datagram has source IP address, destination IP address
 - ▶ each datagram carries 1 transport-layer segment
 - ▶ each segment has source, destination port number
- host uses IP addresses & port numbers to direct segment to appropriate socket



TCP/UDP segment format

Connectionless demultiplexing

- Create sockets with port numbers:

```
addr1.sin_port = htons(12534);
```

```
addr2.sin_port = htons(12535);
```


- UDP socket identified by **two-tuple**:

(dest IP address, dest port number)

- When host receives UDP segment:
 - checks destination port number in segment
 - directs UDP segment to socket with that port number
- IP datagrams with different source IP addresses and/or source port numbers directed to **same socket**



Connectionless demux (cont)




SP provides “return address”


Connection-oriented demux

- TCP socket identified by 4-tuple:
 - ▶ source IP address
 - ▶ source port number
 - ▶ dest IP address
 - ▶ dest port number
- recv host uses all four values to direct segment to appropriate socket
 - Server host may support many simultaneous TCP sockets:
 - ▶ each socket identified by its own 4-tuple
 - Web servers have different sockets for each connecting client
 - ▶ non-persistent HTTP will have different socket for each request

Connection-oriented demux (cont)



Connection-oriented demux: Threaded Web Server



Chapter 3 Outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - ▶ segment structure
 - ▶ reliable data transfer
 - ▶ flow control
 - ▶ connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

UDP: User Datagram Protocol [RFC 768]


- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
 - ▶ lost
 - ▶ delivered out of order to app
- **connectionless:**
 - ▶ no handshaking between UDP sender, receiver
 - ▶ each UDP segment handled independently of others

Why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small segment header
- no congestion control: UDP can blast away as fast as desired

UDP: more

- often used for streaming multimedia apps
 - loss tolerant
 - rate sensitive
- other UDP uses
 - DNS
 - SNMP
- reliable transfer over UDP: add reliability at application layer
 - application-specific error recovery!



UDP segment format

UDP checksum

UDP pseudo-header for checksum computation (IPv4)

Offset	Octet	0								1								2								3																																			
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31																												
0	0	Source Address																																																											
4	32	Destination Address																																																											
8	64	Zeroes				Protocol								UDP Length																																															
12	96	Source Port												Destination Port																																															
16	128	Length												Checksum																																															
20	160	Data																																																											
24	192																																																												
:	:																																																												

Note: This is *not* the actual header that is sent!

UDP checksum

UDP pseudo-header for checksum computation (IPv6)

Offset	Octet	0								1								2								3											
		Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31		
0	0																																				
4	32																																				
8	64																																				
12	96																																				
16	128																																				
20	160																																				
24	192																																				
28	224																																				
32	256																																				
36	288																																				
40	320																																				
44	352																																				
48	384																																				
52	416																																				
:	:																																				

UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

Sender:

- treat segment **contents** as sequence of 16-bit integers
- What exactly is “contents”?
 - Pseudo header taken from selected fields in IP header (included in case packet is misrouted to wrong IP address due to bit error)
 - UDP header
 - UDP data

Receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected. But maybe errors nonetheless?
More later

UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

Sender:

- treat segment **contents** as sequence of 16-bit integers
- What exactly is “contents”?
 - Pseudo header taken from selected fields in IP header (included in case packet is misrouted to wrong IP address due to bit error)
 - UDP header
 - UDP data

Receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected. But maybe errors nonetheless?
More later

UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

Sender:

- treat segment contents as sequence of 16-bit integers
- checksum: addition (1's complement sum) of segment contents
 - How is this different than 2's complement?
- sender puts checksum value into UDP checksum field

Receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected. But maybe errors nonetheless?
More later

Internet Checksum Example


- Note
 - When adding numbers, a carryout from the most significant bit needs to be added to the result
- Example: add two 16-bit integers



	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
	1	0	1	1	1	0	1	1	1	0	1	1	1	0	0	0
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Port Scanning

- Technique used by black- and white-hat communities alike.
- Attempts to connect to a large number (usually all) of ports on a machine.
 - ▶ Successful responses mean that a process is running.
 - ▶ If you know what processes are running, you may be able to select the right exploit to launch.
 - ▶ Most firewalls offer some protection against this.
- This is happening *all the time* on the Internet.
 - ▶ The bad guys are constantly looking for a way in...



Port Scanning Tools

- *nmap* is a popular tool for port scanning.
 - ▶ ...and it is free...
- By seeing which ports are active, nmap can tell a lot about your machine.
 - ▶ For instance, what OS you are running...
- Be careful to check with admins before running this!
 - ▶ Most admins will automatically shut you down if you run it...

The screenshot shows a terminal window with two distinct Nmap scan outputs. The top section is for the host 'scanme.nmap.org' (IP 205.217.153.62) and the bottom section is for the host 'd0ze.internal' (IP 192.168.12.3). Both scans were run with the command '# nmap -A -T4 scanme.nmap.org d0ze'.

Scan Results for scanme.nmap.org:

```
# nmap -A -T4 scanme.nmap.org d0ze
Starting Nmap 4.01 ( http://www.insecure.org/nmap/ ) at 2006-03-20 15:53 PST
Interesting ports on scanme.nmap.org (205.217.153.62):
(The 1667 ports scanned but not shown below are in state: filtered)
PORT      STATE SERVICE VERSION
22/tcp    open  ssh      OpenSSH 3.9p1 (protocol 1.99)
25/tcp    open  smtp     Postfix smtpd
53/tcp    open  domain   ISC Bind 9.2.1
70/tcp    closed gopher
80/tcp    open  http     Apache httpd 2.0.52 ((Fedora))
113/tcp   closed auth
Device type: general purpose
Running: Linux 2.6.X
OS details: Linux 2.6.0 - 2.6.11
Uptime 26.177 days (since Wed Feb 22 11:39:16 2006)

Interesting ports on d0ze.internal (192.168.12.3):
(The 1664 ports scanned but not shown below are in state: closed)
PORT      STATE SERVICE VERSION
21/tcp    open  ftp      Serv-U ftptd 4.0
25/tcp    open  smtp     IMail NT-ESMTP 7.15 2015-2
80/tcp    open  http     Microsoft IIS webserver 5.0
110/tcp   open  pop3    IMail pop3d 7.15 931-1
135/tcp   open  mtask   Microsoft mtask (task server - c:\winnt\system32\
139/tcp   open  netbios-ssn
445/tcp   open  microsoft-ds Microsoft Windows XP microsoft-ds
1025/tcp  open  msrpc   Microsoft Windows RPC
5800/tcp  open  vnc-http Ultr@VNC (Resolution 1024x800; VNC TCP port: 5900)
MAC Address: 00:A0:CC:51:72:7E (Lite-on Communications)
Device type: general purpose
Running: Microsoft Windows NT/2K/XP
OS details: Microsoft Windows 2000 Professional
Service Info: OS: Windows

Nmap finished: 2 IP addresses (2 hosts up) scanned in 42.291 seconds
```