

CMSC 332

Computer Networks

Reliable Data Transfer

Professor Szajda

Last Time

- Multiplexing/Demultiplexing at the Transport Layer.
 - ▶ How do TCP and UDP differ?
- UDP gives us virtually “bare-bones” access to the network layer.
 - ▶ What are the four fields in a UDP header?
- What is port scanning?
 - ▶ How can it be used to protect systems? To attack them?




Chapter 3 outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - ▶ segment structure
 - ▶ reliable data transfer
 - ▶ flow control
 - ▶ connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control


Problem Solving

- Given all of our talk about TCP, you might assume that describing how it works is straightforward.
- The reality is that there are lots of different ways that we could have provided “reliable” delivery.
 - ▶ Here is where we will really start to see design tradeoffs.
- It’s time in your CS career to be creative...
 - ▶ Let’s solve some problems.



Principles of Reliable data transfer

- important in application, transport, link layers
- top-10 list of important networking topics!




- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Reliable data transfer: getting started


rdt_send() : called from above, (e.g., by app.). Passed data to deliver to receiver upper layer

send
side



deliver_data() : called by rdt to deliver data to upper layer

receive
side



udt_send() : called by rdt, to transfer packet over unreliable channel to receiver


rdt_rcv() : called when packet arrives on rcv-side of channel

Reliable data transfer: getting started

We'll:


- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
 - but control info will flow in **both** directions!
- use finite state machines (FSM) to specify sender, receiver

state: when in this “state” next state uniquely determined by next event



Rdt1.0: reliable transfer over a reliable channel


- underlying channel perfectly reliable
 - no bit errors
 - no loss of packets
- separate FSMs for sender, receiver:
 - sender sends data into underlying channel
 - receiver read data from underlying channel




Rdt2.0: channel with bit errors

- underlying channel may flip bits in packet
 - ▶ checksum to detect bit errors
- *the question: how to recover from errors:*
 - ▶ **acknowledgements (ACKs):** receiver explicitly tells sender that pkt received OK
 - ▶ **negative acknowledgements (NAKs):** receiver explicitly tells sender that pkt had errors
 - ▶ sender retransmits pkt on receipt of NAK
- new mechanisms in `rdt2 . 0` (beyond `rdt1 . 0`):
 - ▶ error detection
 - ▶ receiver feedback: control msgs (ACK,NAK) rcvr->sender


rdt2.0: FSM specification




rdt2.0: operation with no errors




rdt2.0: error scenario



rdt2.0: error scenario



rdt2.0: error scenario



No. Why? And how do we fix things?


rdt2.0 has a fatal flaw!

What happens if ACK/NAK corrupted? (but does arrive)

- sender doesn't know what happened at receiver!
- can't just retransmit: possible duplicate

Handling duplicates:


- sender retransmits current pkt if ACK/NAK garbled
- sender adds sequence number to each pkt
- receiver discards (doesn't deliver up) duplicate pkt




stop and wait

Sender sends one packet,
then waits for receiver
response

rdt2.1: sender, handles garbled ACK/NAKs



rdt2.1: receiver, handles garbled ACK/NAKs



rdt2.1: discussion

Sender:

- seq # added to pkt
- two seq. #'s (0,1) will suffice.
Why?
- must check if received ACK/NAK corrupted
- twice as many states
 - state must “remember” whether “current” pkt has 0 or 1 seq. #

Receiver:


- must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #
- note: receiver can not know if its last ACK/NAK received OK at sender

rdt2.2: a NAK-free protocol

- same functionality as rdt2.1, using ACKs only
- instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must **explicitly** include seq # of pkt being ACKed
- duplicate ACK at sender results in same action as NAK: *retransmit current pkt*



rdt2.2: sender, receiver fragments



rdt3.0: channels with errors and loss

New assumption: underlying channel can also lose packets (data or ACKs)


- checksum, seq. #, ACKs, retransmissions will be of help, but not enough




Approach: sender waits “reasonable” amount of time for ACK

- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but use of seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed
- requires countdown timer


rdt3.0 sender



rdt3.0 in action




(a) operation with no loss




(b) lost packet

rdt3.0 in action



(c) lost ACK



(d) premature timeout

Performance of rdt3.0

- rdt3.0 works, but performance very poor
- example: 1 Gbps link, 15 ms e-e prop. delay, 1 KB packet:
 - ▶ U_{sender} : **utilization** – fraction of time sender busy sending

Performance of rdt3.0

- rdt3.0 works, but performance very poor
- example: 1 Gbps link, 15 ms e-e prop. delay, 1 KB packet:


$$T_{\text{transmit}} = \frac{L \text{ (packet length in bits)}}{R \text{ (transmission rate, bps)}} = \frac{8\text{kb/pkt}}{10^9 \text{ b/sec}} = 8 \text{ microsec}$$

- ▶ U_{sender} : **utilization** – fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$


- ▶ 1KB pkt every 30 msec -> 33kB/sec throughput over 1 Gbps link
- ▶ network protocol limits use of physical resources!

rdt3.0: stop-and-wait operation



$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

rdt3.0: stop-and-wait operation




So, question: How do we improve this?

Pipelined protocols


Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- ▶ range of sequence numbers must be increased
- ▶ buffering at sender and/or receiver



- Two generic forms of pipelined protocols:
go-Back-N, selective repeat

Pipelining: increased utilization



Pipelining Protocols

- Go-back-N: big picture:
 - ▶ Sender can have up to N unacked packets in pipeline
 - ▶ Rcvr only sends cumulative acks
 - Doesn't ack packet if there's a gap
 - ▶ Sender has timer for oldest unacked packet
 - If timer expires, retransmit all unacked packets



Selective Repeat: Big Picture


- Sender can have up to N unacked packets in pipeline
- Rcvr acks individual packets
- Sender maintains timer for each unacked packet
 - ▶ When timer expires, retransmit only individual unacked packet.



Go-Back-N


Sender:

- k-bit seq # in pkt header
- “window” of up to N, consecutive unack’ed pkts allowed




- ACK(n):ACKs all pkts up to, including seq # n - “cumulative ACK”
 - ▶ may receive duplicate ACKs (see receiver)
- timer for each in-flight pkt
- timeout(n): retransmit pkt n and all higher seq # pkts in window

GBN: sender extended FSM




GBN: receiver extended FSM



ACK-only: always send ACK for correctly-received pkt with highest **in-order** seq #

- may generate duplicate ACKs
- need only remember **expectedseqnum**
- out-of-order pkt:
 - discard (don't buffer) -> **no receiver buffering!**
 - Re-ACK pkt with highest in-order seq #


GBN in action



Selective Repeat

- receiver *individually* acknowledges all correctly received pkts
 - ▶ buffers pkts, as needed, for eventual in-order delivery to upper layer
- sender only resends pkts for which ACK not received
 - ▶ sender timer for each unACKed pkt
- sender window
 - ▶ N consecutive seq #'s
 - ▶ again limits seq #'s of sent, unACKed pkts

Selective repeat: sender, receiver windows



Selective repeat

sender

data from above :

- if next available seq # in window, send pkt

timeout(n):

- resend pkt n, restart timer

ACK(n) in [sendbase,sendbase+N]:

- mark pkt n as received
- if n smallest unACKed pkt, advance window base to next unACKed seq #

receiver

pkt n in [rcvbase, rcvbase+N-1]

- send ACK(n)
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt


pkt n in [rcvbase-N,rcvbase-1]

- ACK(n)

otherwise:



- ignore

Selective repeat in action



Selective repeat: dilemma



What is going on here?
Put another way, what
is this slide demonstrating?
How do we fix it?



Selective repeat: dilemma

Example:

- seq #'s: 0, 1, 2, 3
- window size=3
- receiver sees no difference in two scenarios!
- incorrectly passes duplicate data as new in (a)



Q: what relationship between seq # size and window size?

Next Time

- That was a lot of material...
 - ▶ Take time to look over the notes. Understand the differences between each of these schemes!
- Next Time
 - ▶ TCP and Congestion Control (Sections 3.5 and 3.6)
- Time for Project 3!
 - ▶ Coming later today!

