

CMSC 332

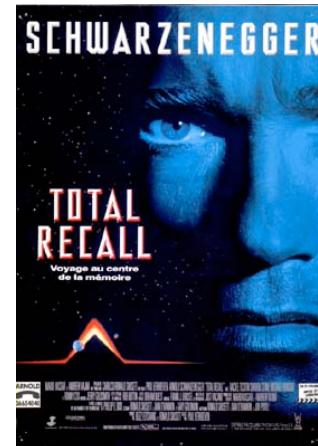
Computer Networks

IP

Professor Szajda

Last Time

- What is forwarding? Routing?
- What's the difference between TCP and VC networks?
- What is longest prefix matching? How does it work?
- What is additive increase/multiplicative decrease?

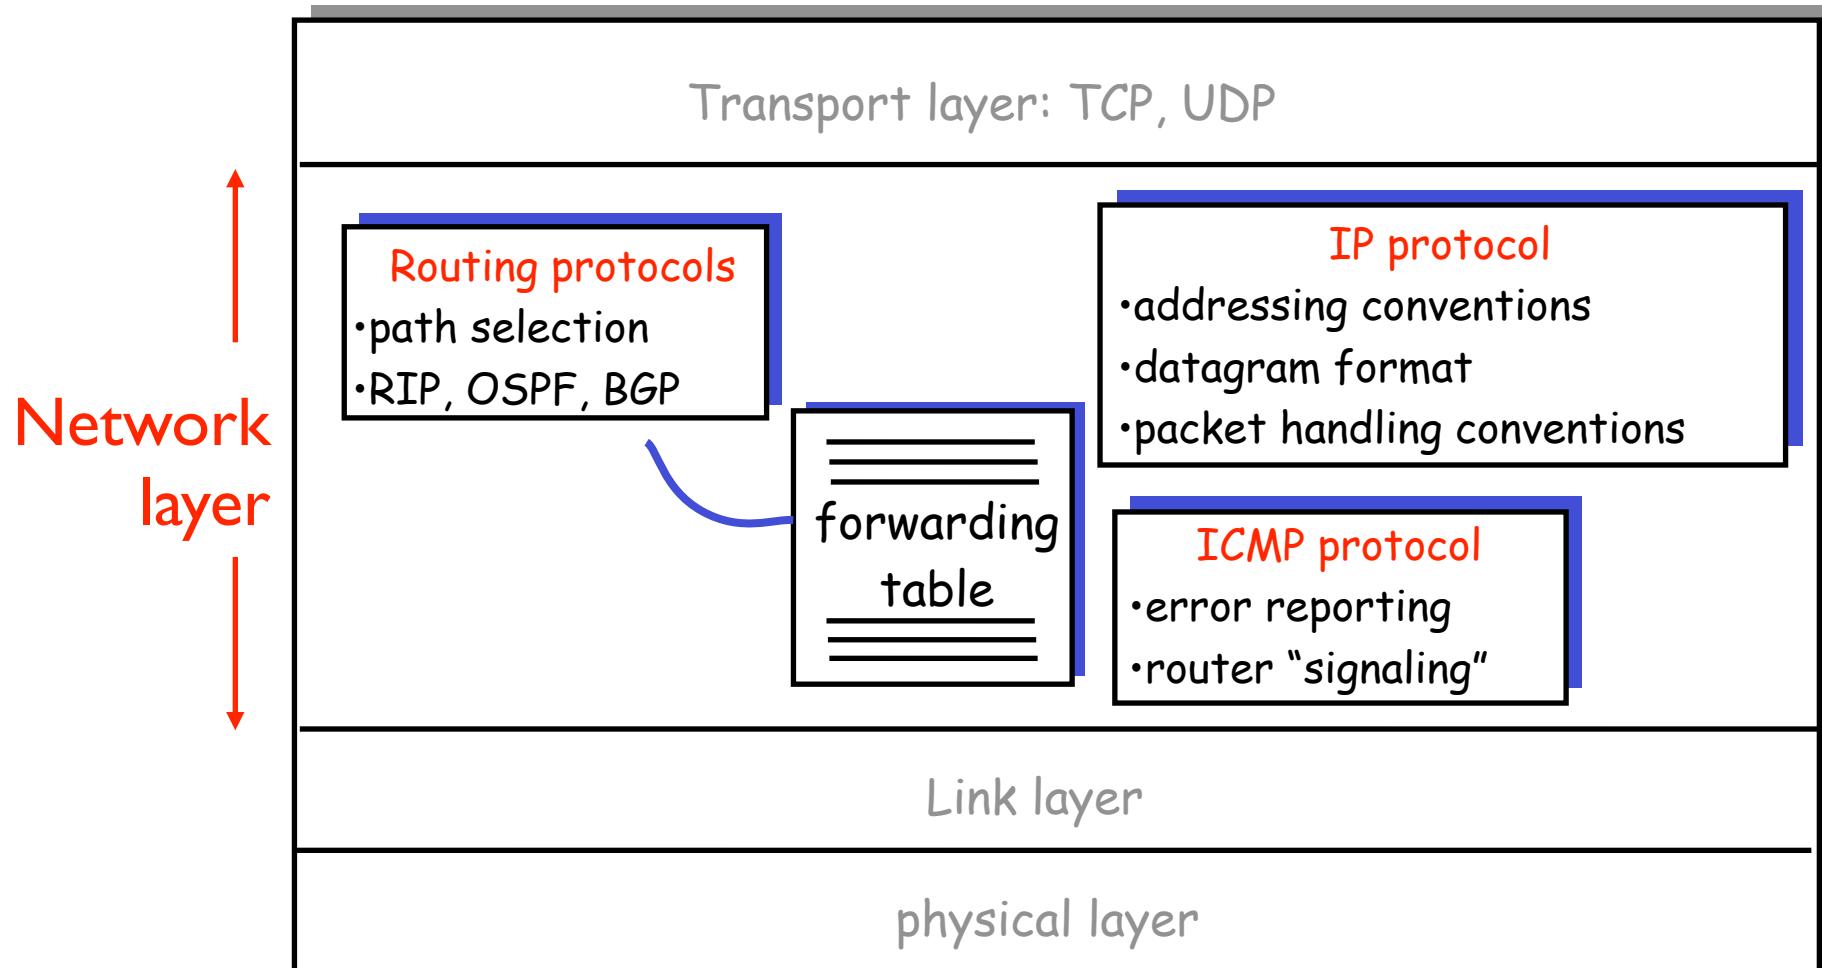


Chapter 4: Network Layer

- 4.1 Introduction
- 4.2 Virtual circuit and datagram networks
- 4.3 What's inside a router
- 4.4 IP: Internet Protocol
 - ▶ Datagram format
 - ▶ IPv4 addressing
 - ▶ ICMP
 - ▶ IPv6
- 4.5 Routing algorithms
 - ▶ Link state
 - ▶ Distance Vector
 - ▶ Hierarchical routing
- 4.6 Routing in the Internet
 - ▶ RIP
 - ▶ OSPF
 - ▶ BGP
- 4.7 Broadcast and multicast routing

The Internet Network layer

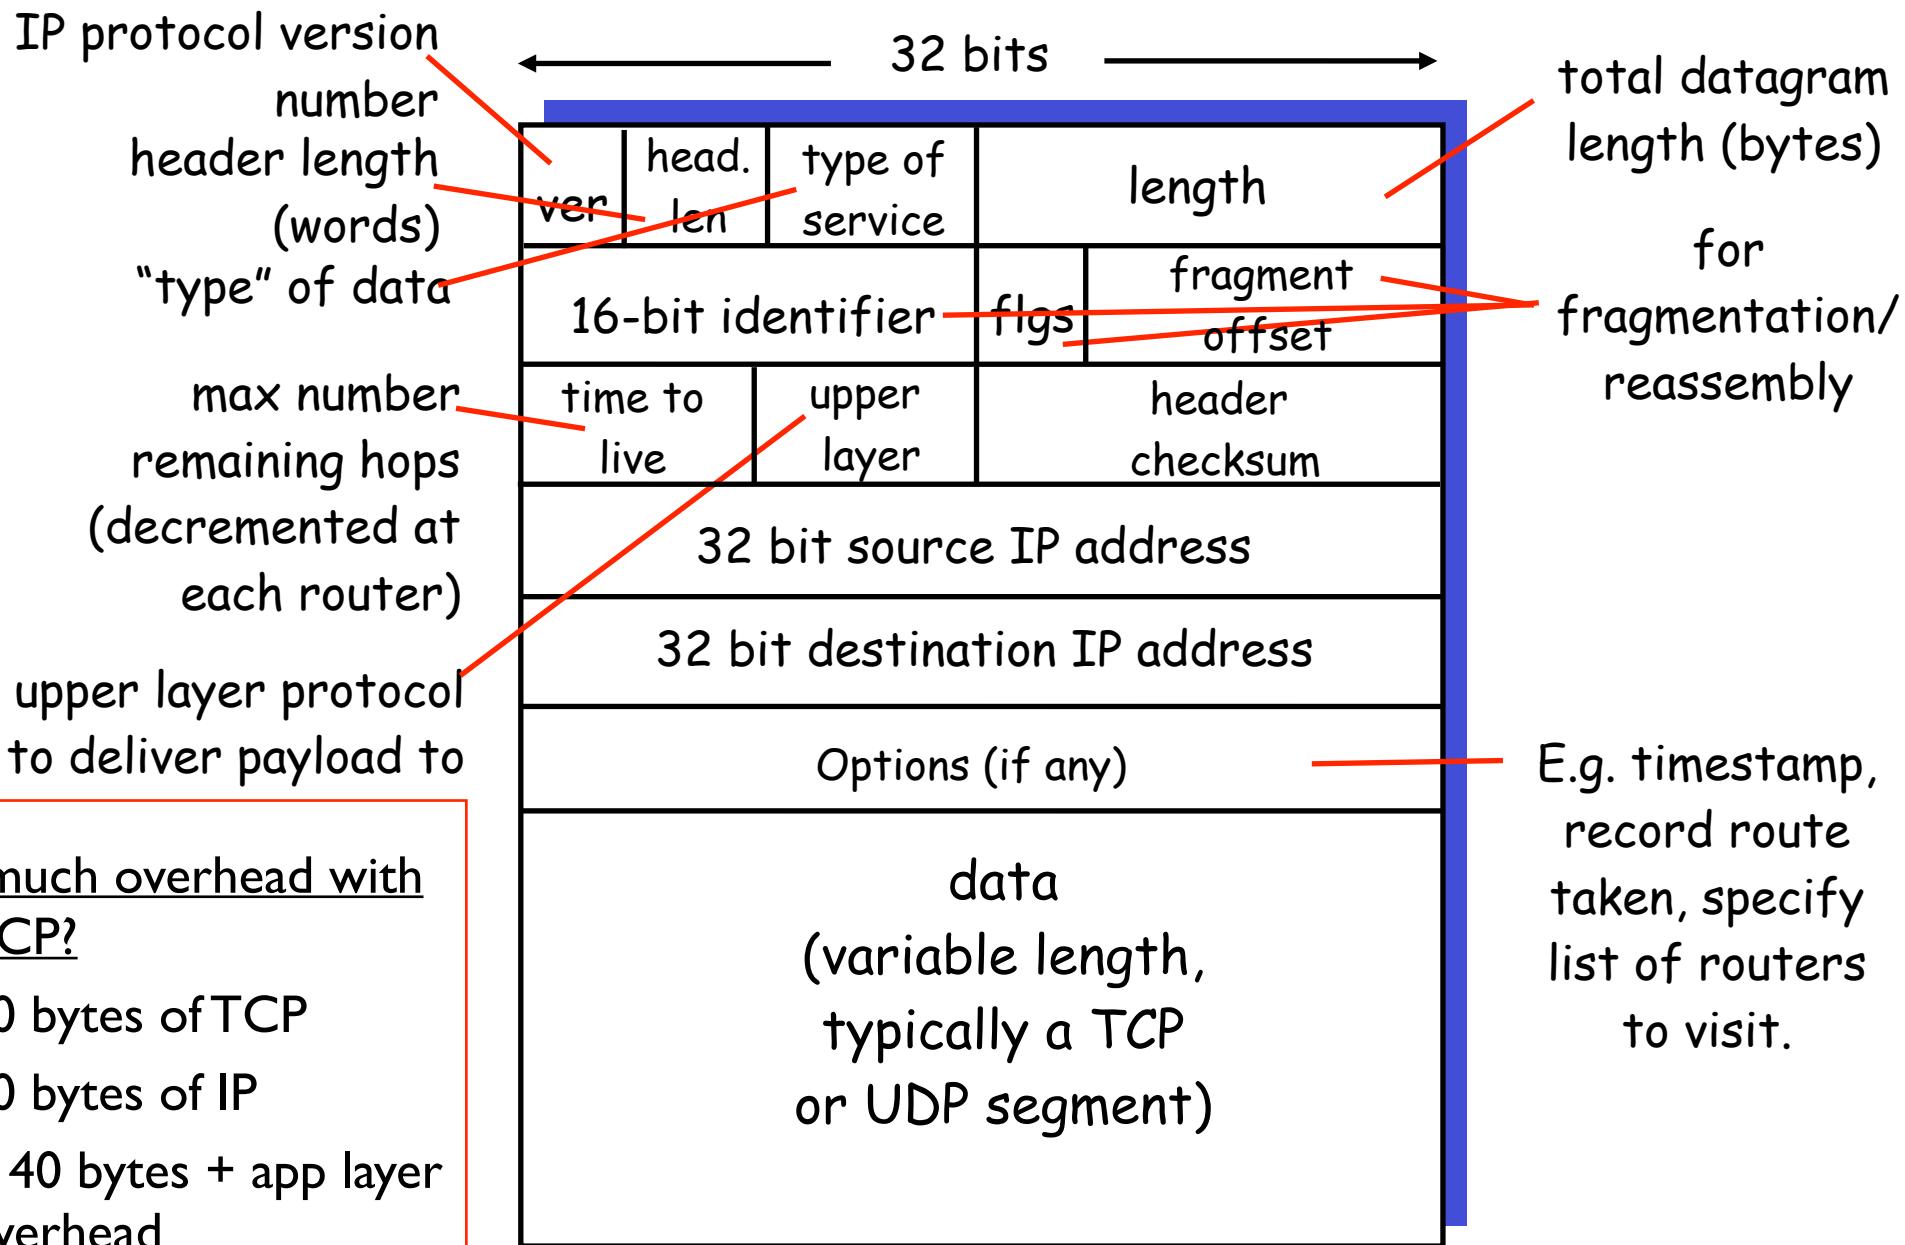
Host, router network layer functions:



Chapter 4: Network Layer

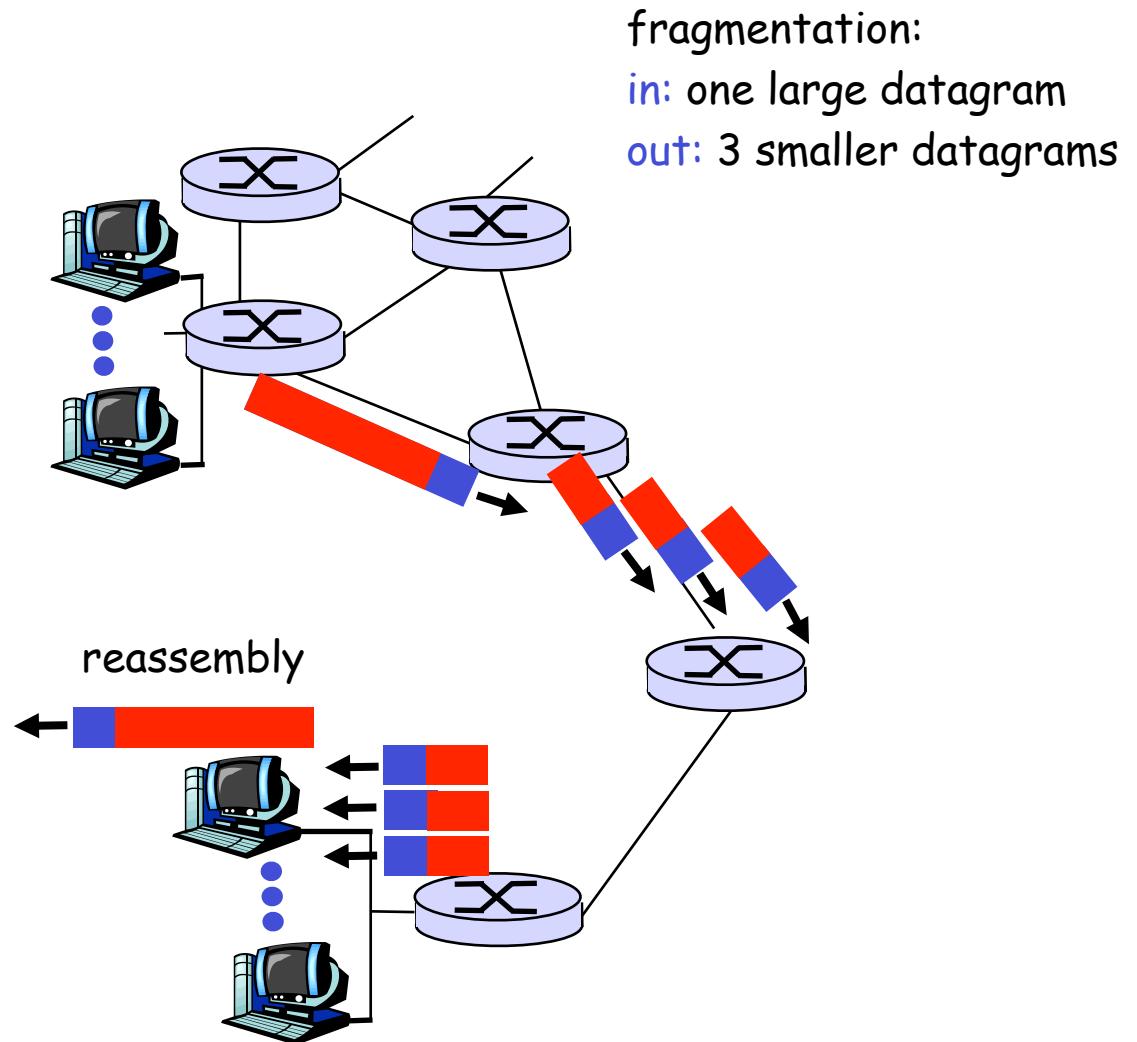
- 4.1 Introduction
- 4.2 Virtual circuit and datagram networks
- 4.3 What's inside a router
- 4.4 IP: Internet Protocol
 - ▶ Datagram format
 - ▶ IPv4 addressing
 - ▶ ICMP
 - ▶ IPv6
- 4.5 Routing algorithms
 - ▶ Link state
 - ▶ Distance Vector
 - ▶ Hierarchical routing
- 4.6 Routing in the Internet
 - ▶ RIP
 - ▶ OSPF
 - ▶ BGP
- 4.7 Broadcast and multicast routing

IP datagram format



IP Fragmentation & Reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame.
 - different link types, different MTUs (Why?)
- large IP datagram divided (“fragmented”) within net
 - one datagram becomes several datagrams
 - “reassembled” only at final destination
 - IP header bits used to identify, order related fragments



IP Fragmentation and Reassembly



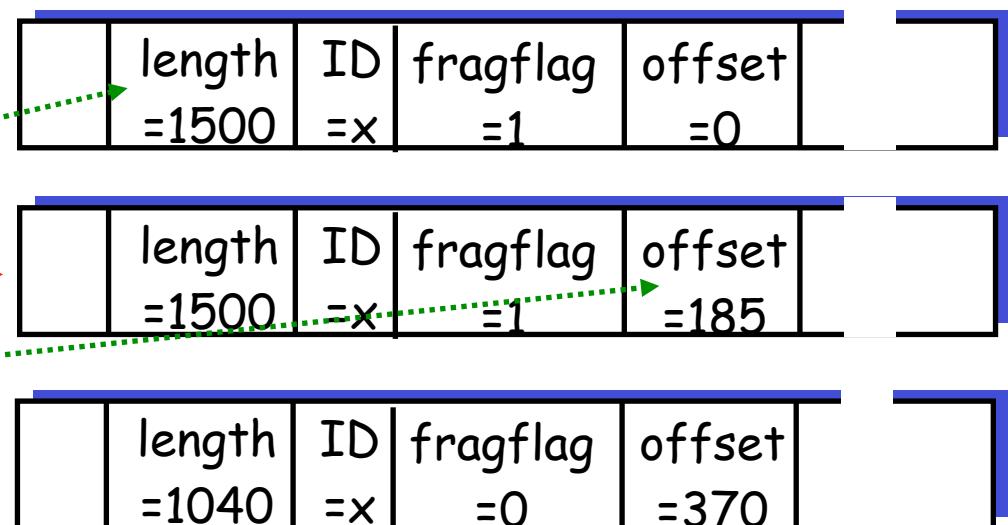
Example

- 4000 byte datagram
- MTU = 1500 bytes

One large datagram becomes several smaller datagrams

1480 bytes in data field

offset =
 $1480/8$

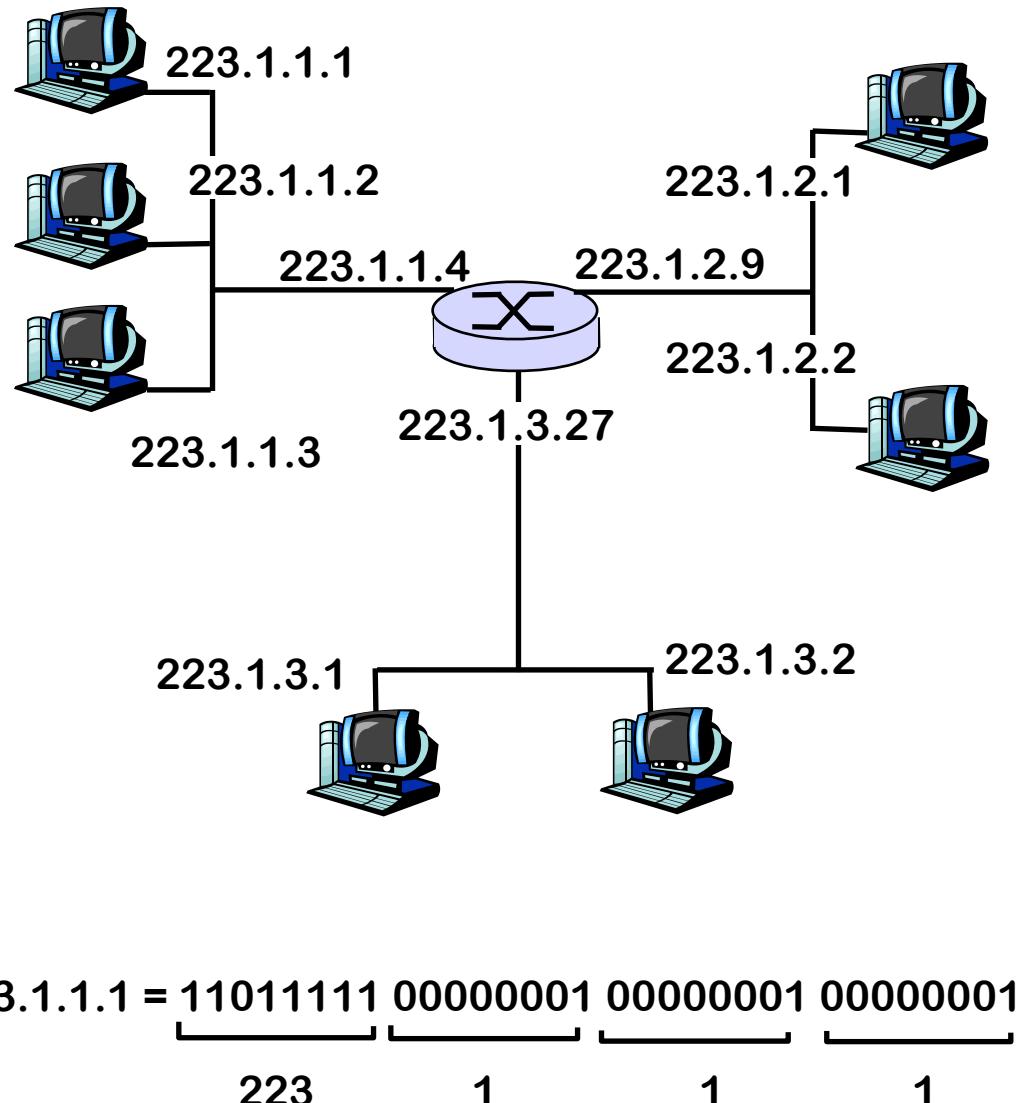


Chapter 4: Network Layer

- 4.1 Introduction
- 4.2 Virtual circuit and datagram networks
- 4.3 What's inside a router
- 4.4 IP: Internet Protocol
 - ▶ Datagram format
 - ▶ IPv4 addressing
 - ▶ ICMP
 - ▶ IPv6
- 4.5 Routing algorithms
 - ▶ Link state
 - ▶ Distance Vector
 - ▶ Hierarchical routing
- 4.6 Routing in the Internet
 - ▶ RIP
 - ▶ OSPF
 - ▶ BGP
- 4.7 Broadcast and multicast routing

IP Addressing: introduction

- **IP address:** 32-bit identifier for host, router interface
- **interface:** connection between host/router and physical link
 - ▶ router's typically have multiple interfaces
 - ▶ host typically has one interface
 - ▶ IP addresses associated with each interface



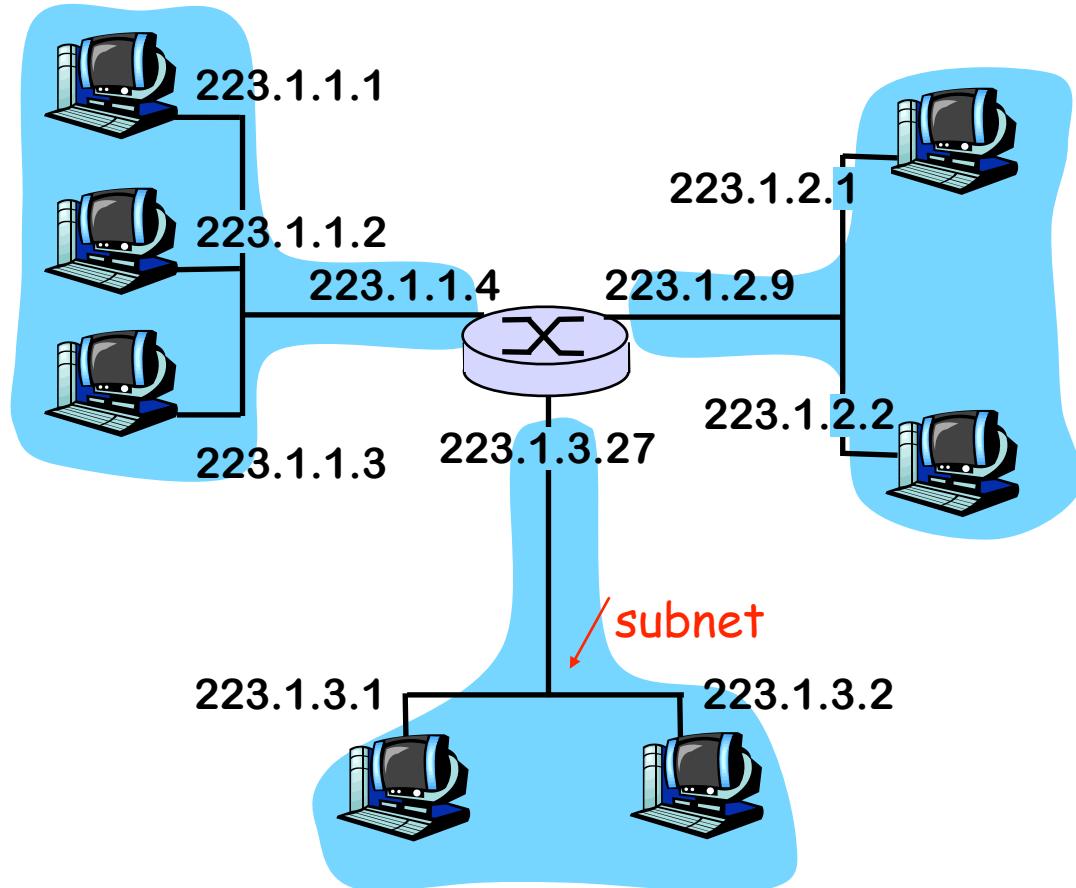
Subnets

- **IP address:**

- subnet part (high order bits)
- host part (low order bits)

- **What's a subnet ?**

- device interfaces with same subnet part of IP address
- can physically reach each other without intervening router

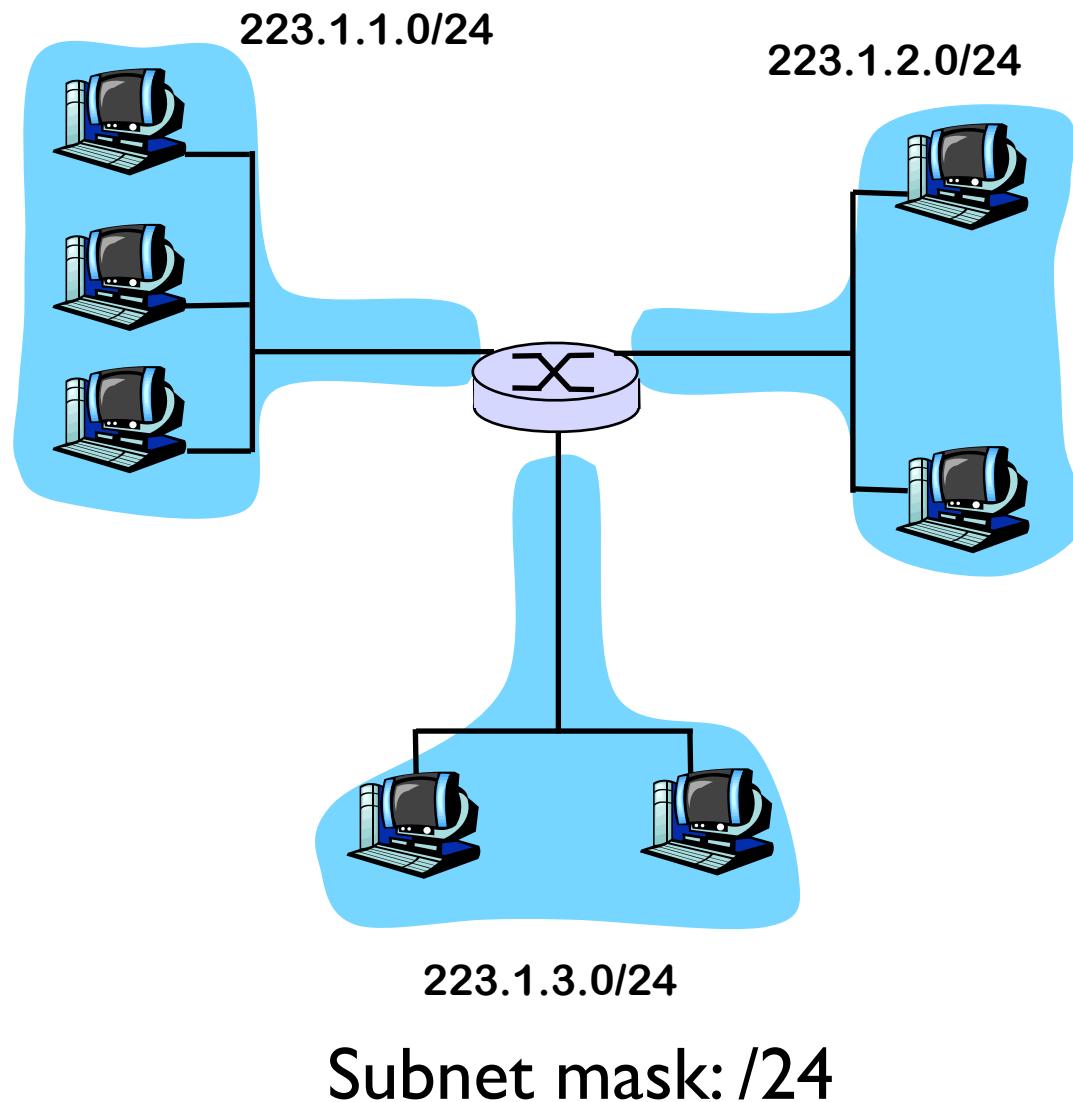


network consisting of 3 subnets

Subnets

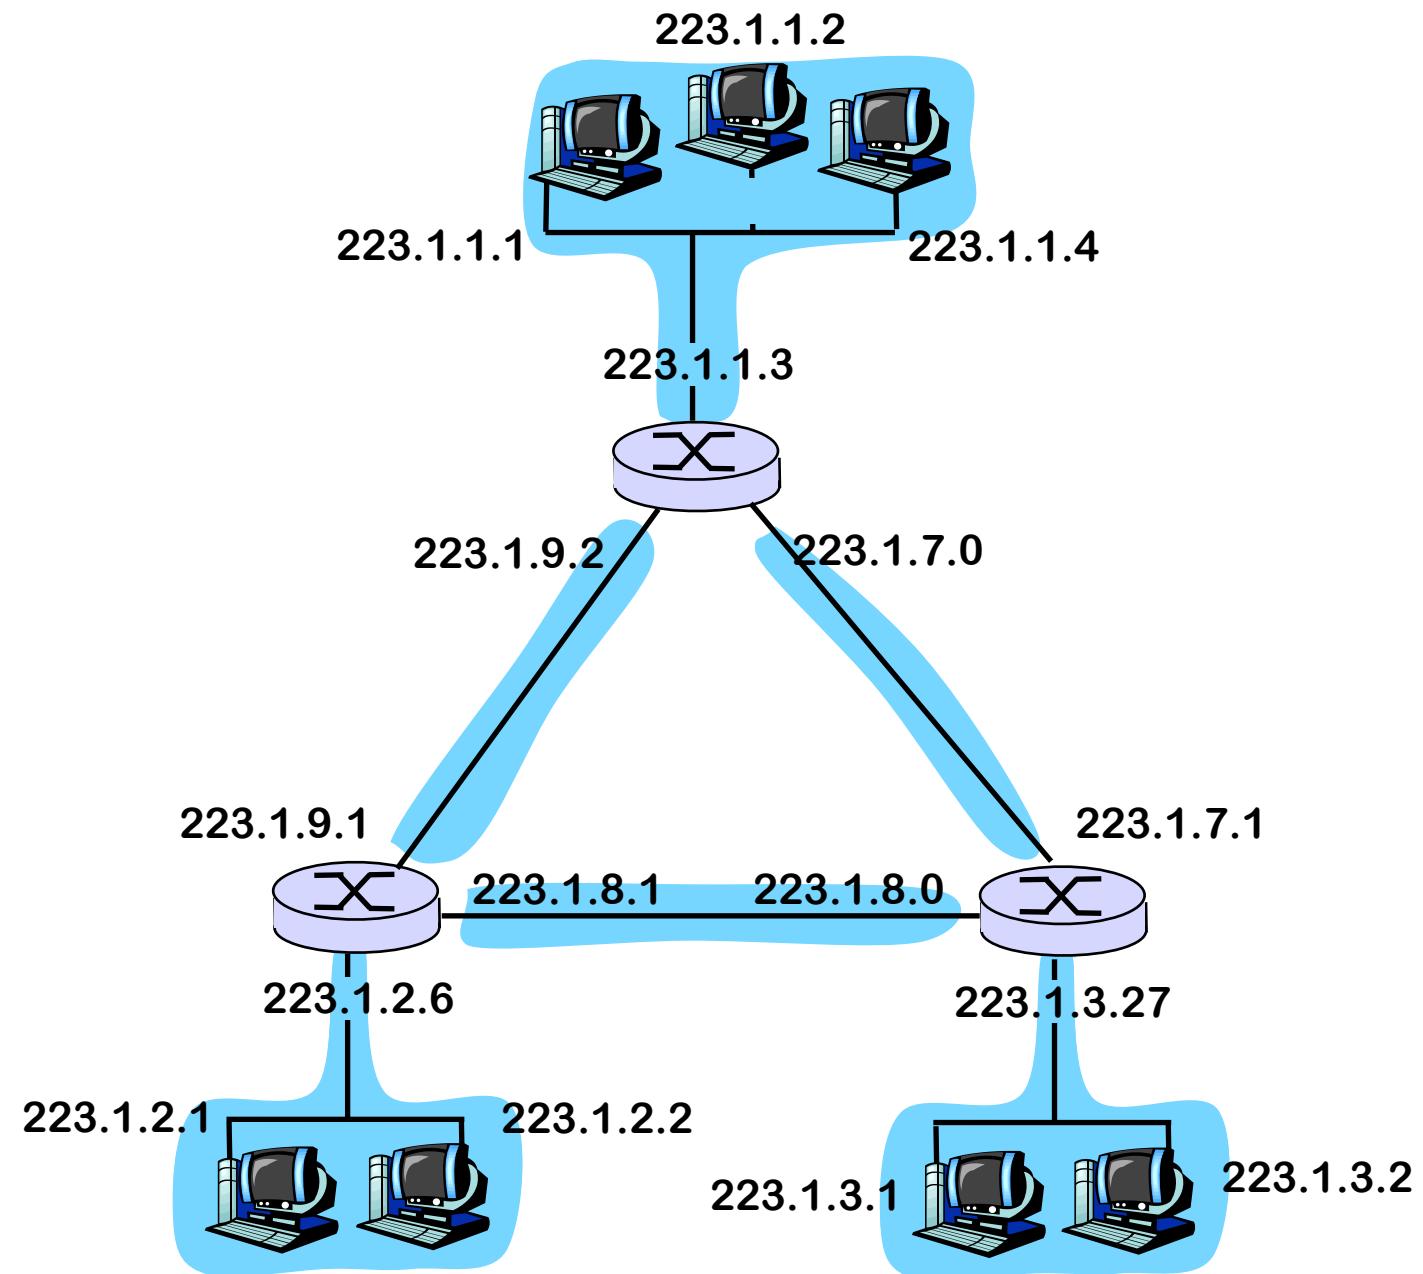
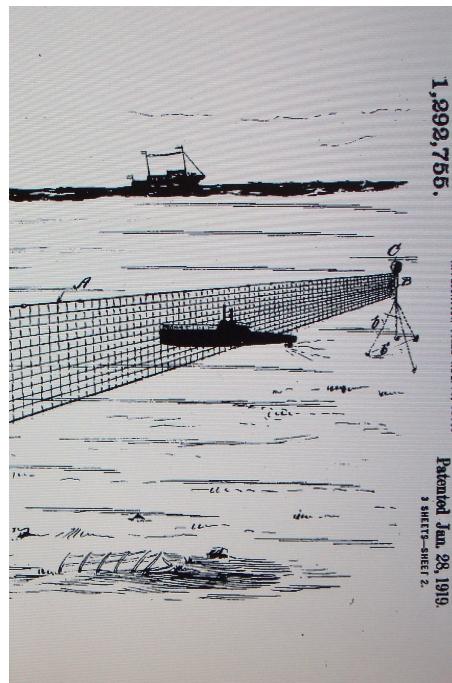
Recipe

- To determine the subnets, detach each interface from its host or router, creating islands of isolated networks. Each isolated network is called a **subnet**.



Subnets

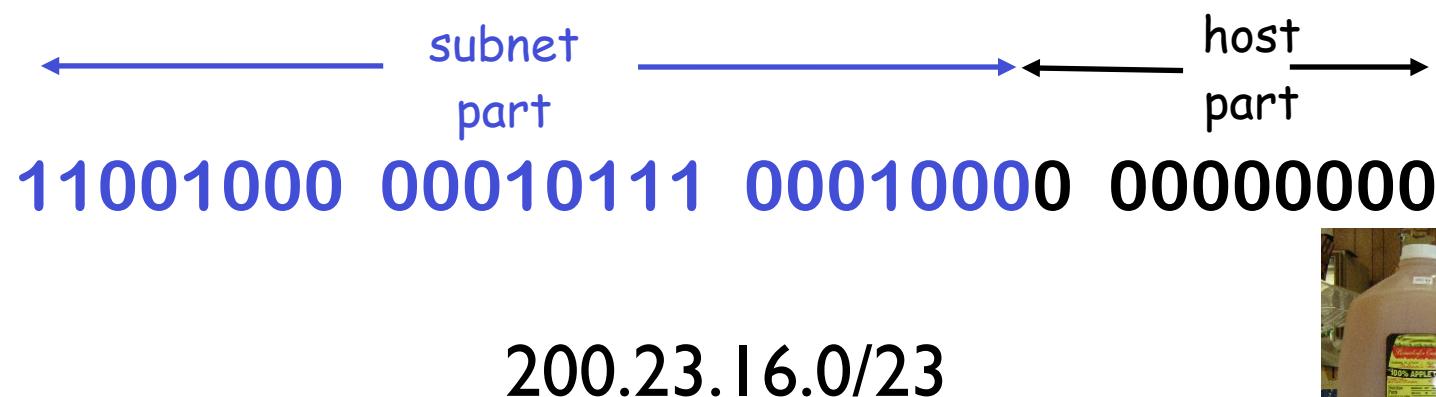
How many?



IP addressing: CIDR

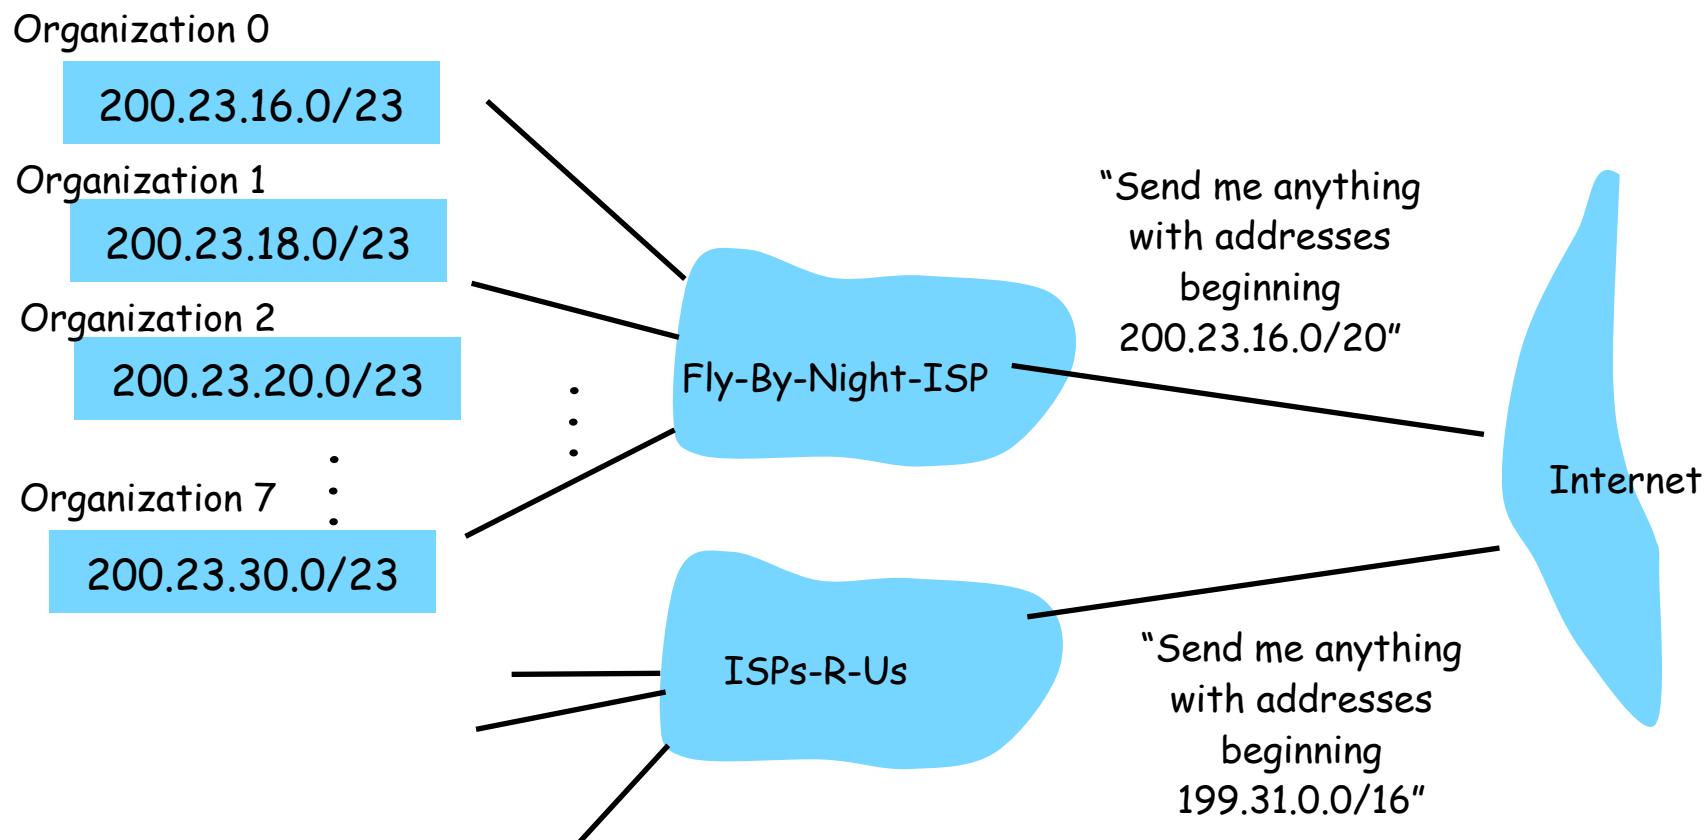
CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
 - address format: $a.b.c.d/x$, where x is # bits in subnet portion of address



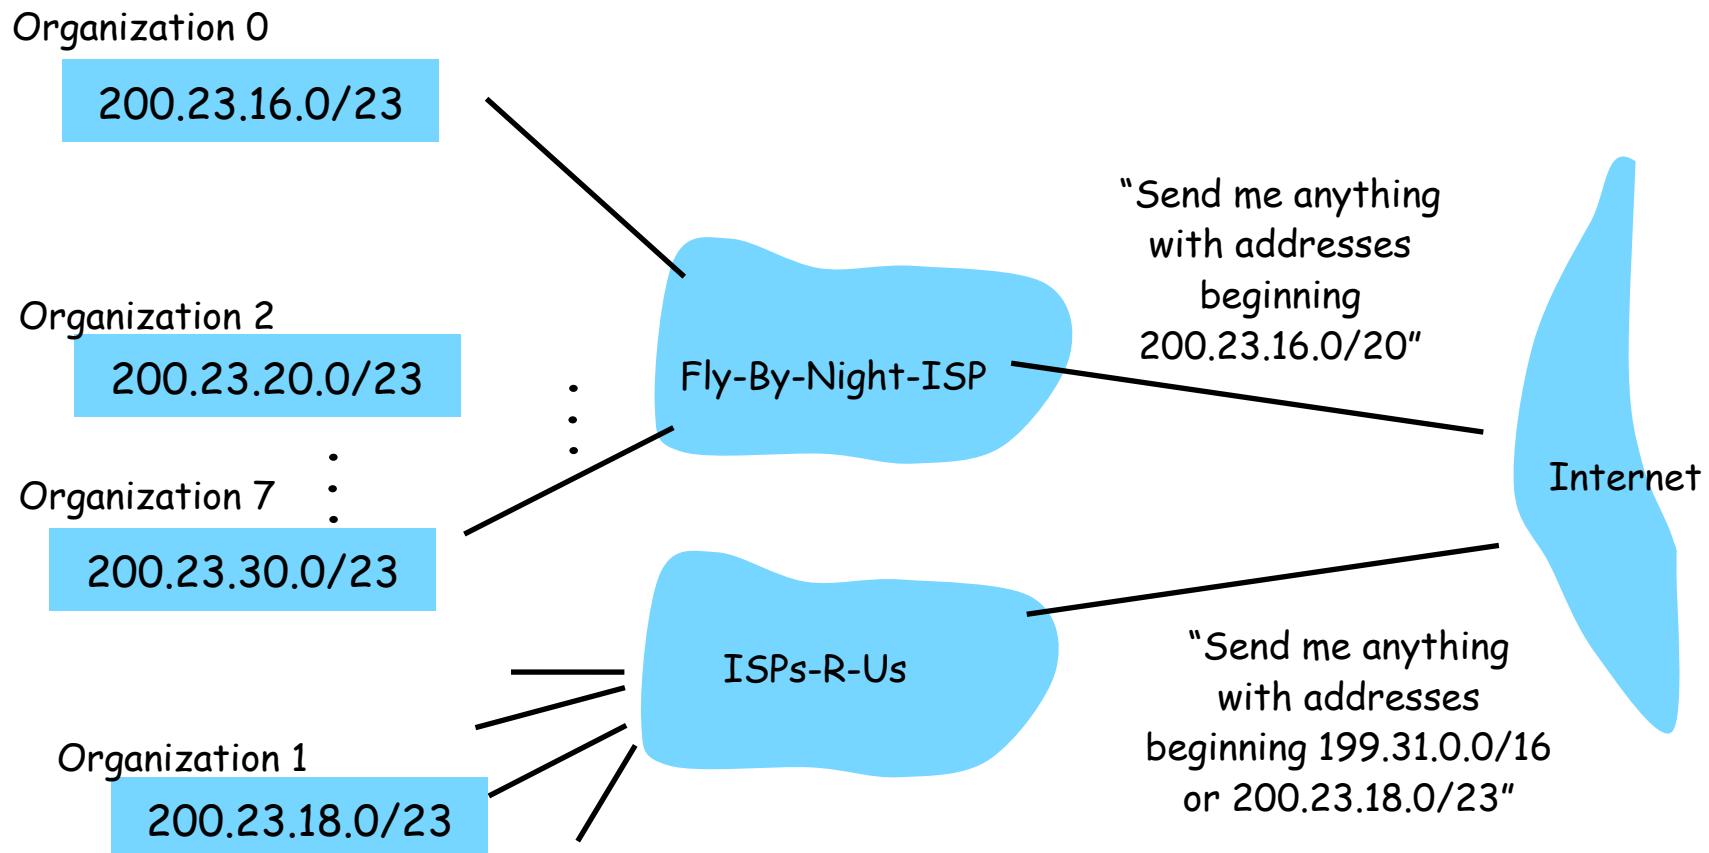
Hierarchical addressing: route aggregation

Hierarchical addressing allows efficient advertisement of routing information:



Hierarchical addressing: more specific routes

ISPs-R-Us has a more specific route to Organization 1



IP addresses: how to get one?

Q: How does a **network** get the subnet part of IP addr?

A: gets allocated portion of its provider ISP's address space

ISP's block 11001000 00010111 00010000 00000000
200.23.16.0/20

Organization 0 11001000 00010111 00010000 00000000
200.23.16.0/23

Organization 1 11001000 00010111 00010010 00000000
200.23.18.0/23

Organization 2 11001000 00010111 00010100 00000000
200.23.20.0/23

...
Organization 7 11001000 00010111 00011110 00000000
200.23.30.0/23

IP addressing: the last word...

Q: How does an ISP get block of addresses?

A: ICANN: Internet Corporation for Assigned Names and Numbers

- ▶ allocates addresses
- ▶ manages DNS
- ▶ assigns domain names, resolves disputes



IP addresses: how to get one?

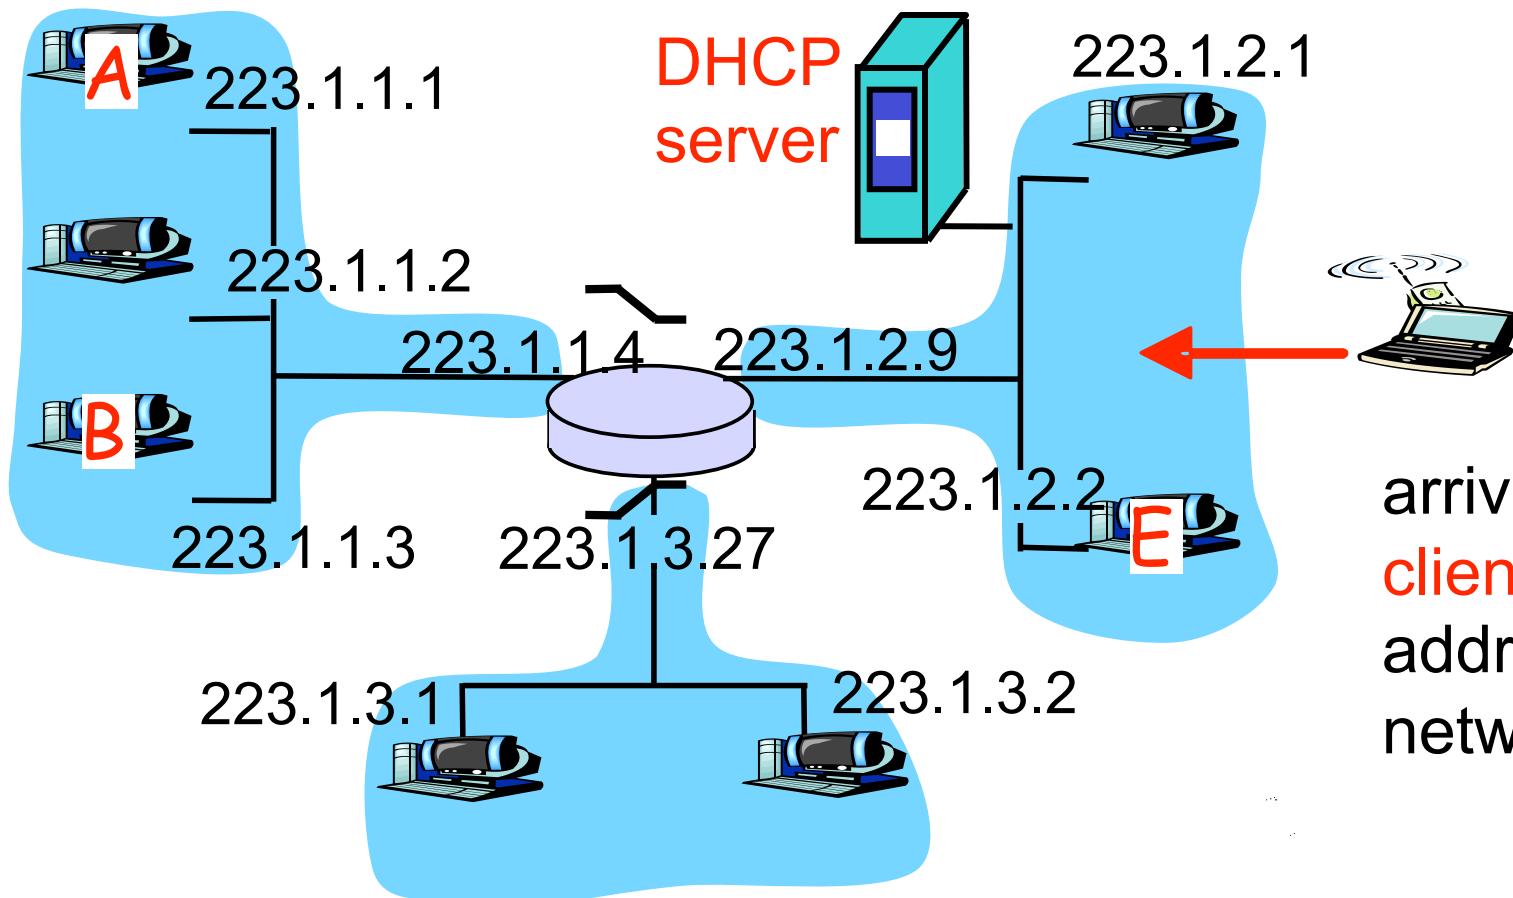
Q: How does a host get an IP address?

- hard-coded by system admin in a file
 - ▶ Windows: control-panel->network->configuration->tcp/ip->properties
 - ▶ UNIX: /etc/rc.config
- **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from a server
 - ▶ “plug-and-play”

DHCP: Dynamic Host Configuration Protocol

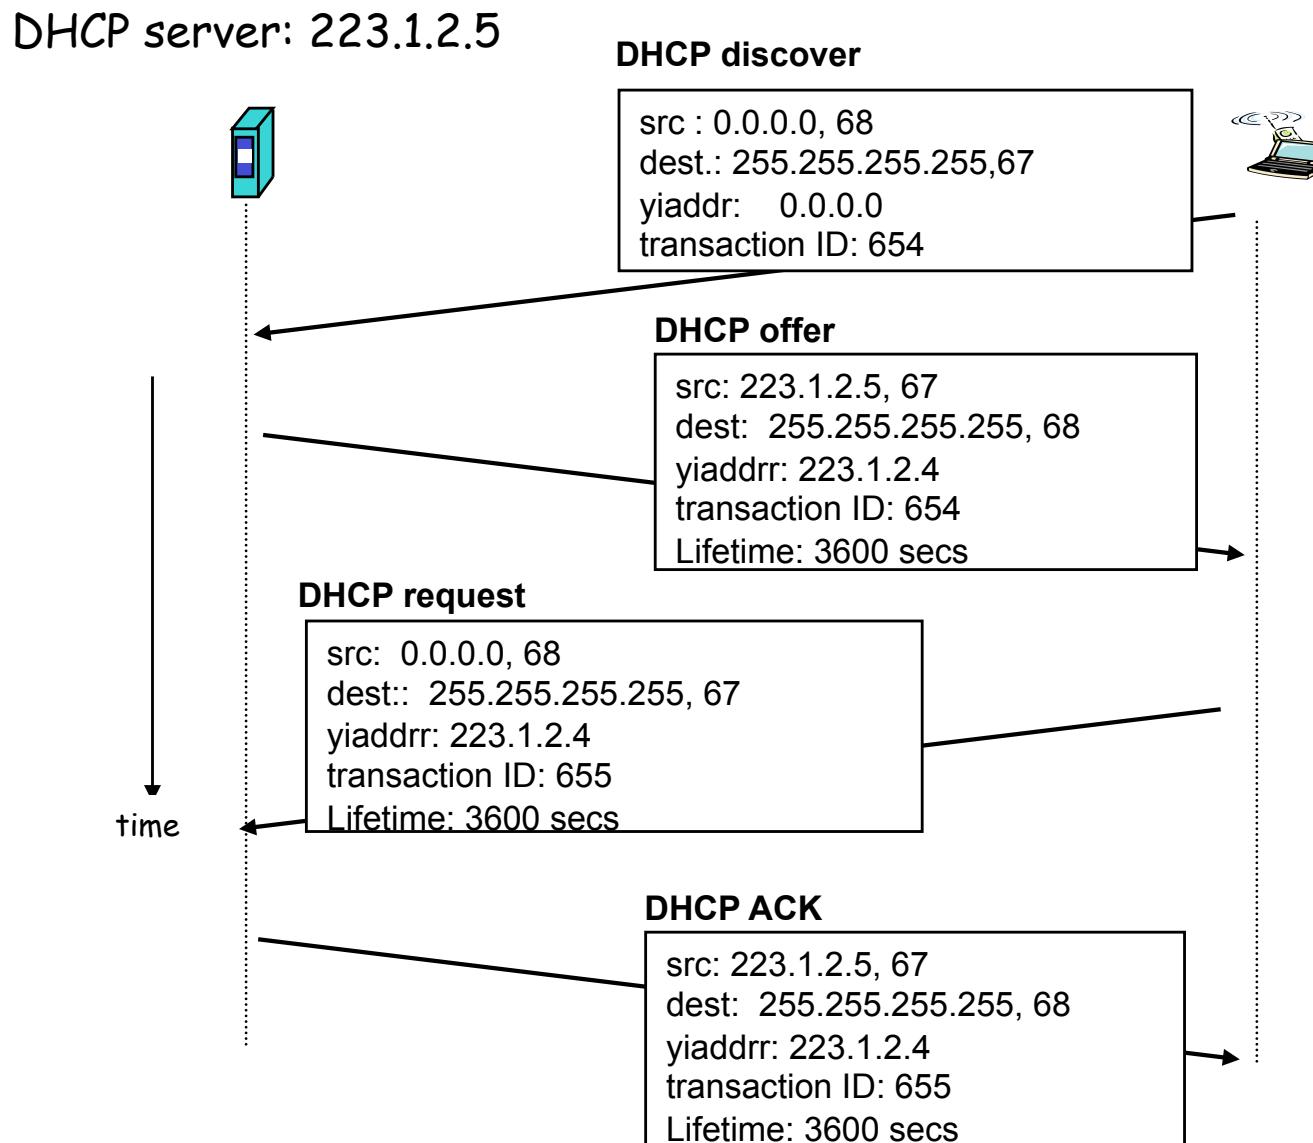
- Goal: allow host to dynamically obtain its IP address from network server when it joins network
 - ▶ Can renew its lease on address in use
 - ▶ Allows reuse of addresses (only hold address while connected and “on”)
 - ▶ Support for mobile users who want to join network (more shortly)
- DHCP overview:
 - ▶ host broadcasts “DHCP discover” msg
 - ▶ DHCP server responds with “DHCP offer” msg
 - ▶ host requests IP address: “DHCP request” msg
 - ▶ DHCP server sends address: “DHCP ack” msg

DHCP client-server scenario

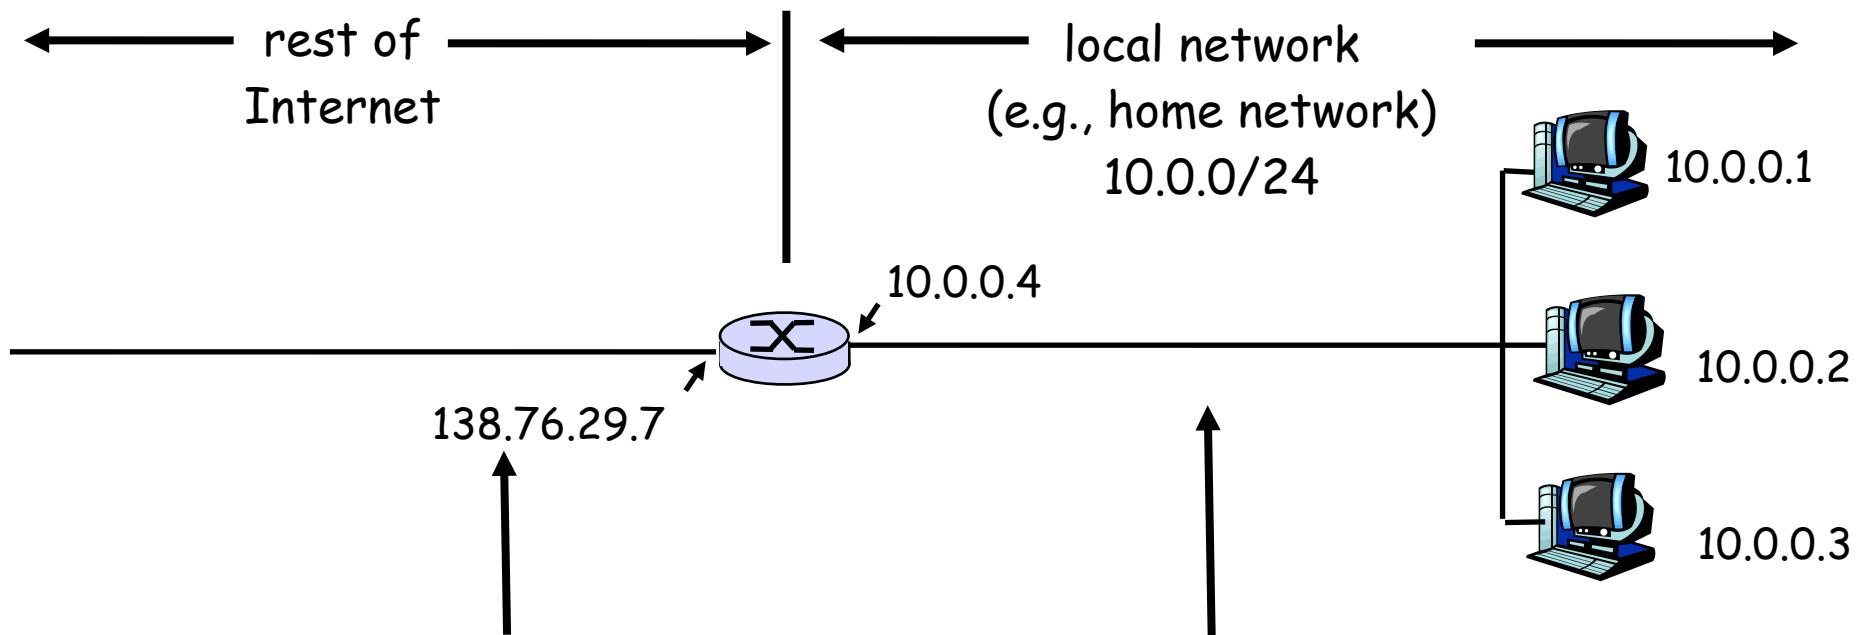


arriving **DHCP**
client needs
address in this
network

DHCP client-server scenario



NAT: Network Address Translation



All datagrams leaving local network have same single source NAT IP address: 138.76.29.7, different source port numbers

Datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

NAT: Network Address Translation

- **Motivation:** local network uses just one IP address as far as outside world is concerned:
 - range of addresses not needed from ISP: just one IP address for all devices
 - can change addresses of devices in local network without notifying outside world
 - can change ISP without changing addresses of devices in local network
 - devices inside local net not explicitly addressable, visible by outside world (a security plus).

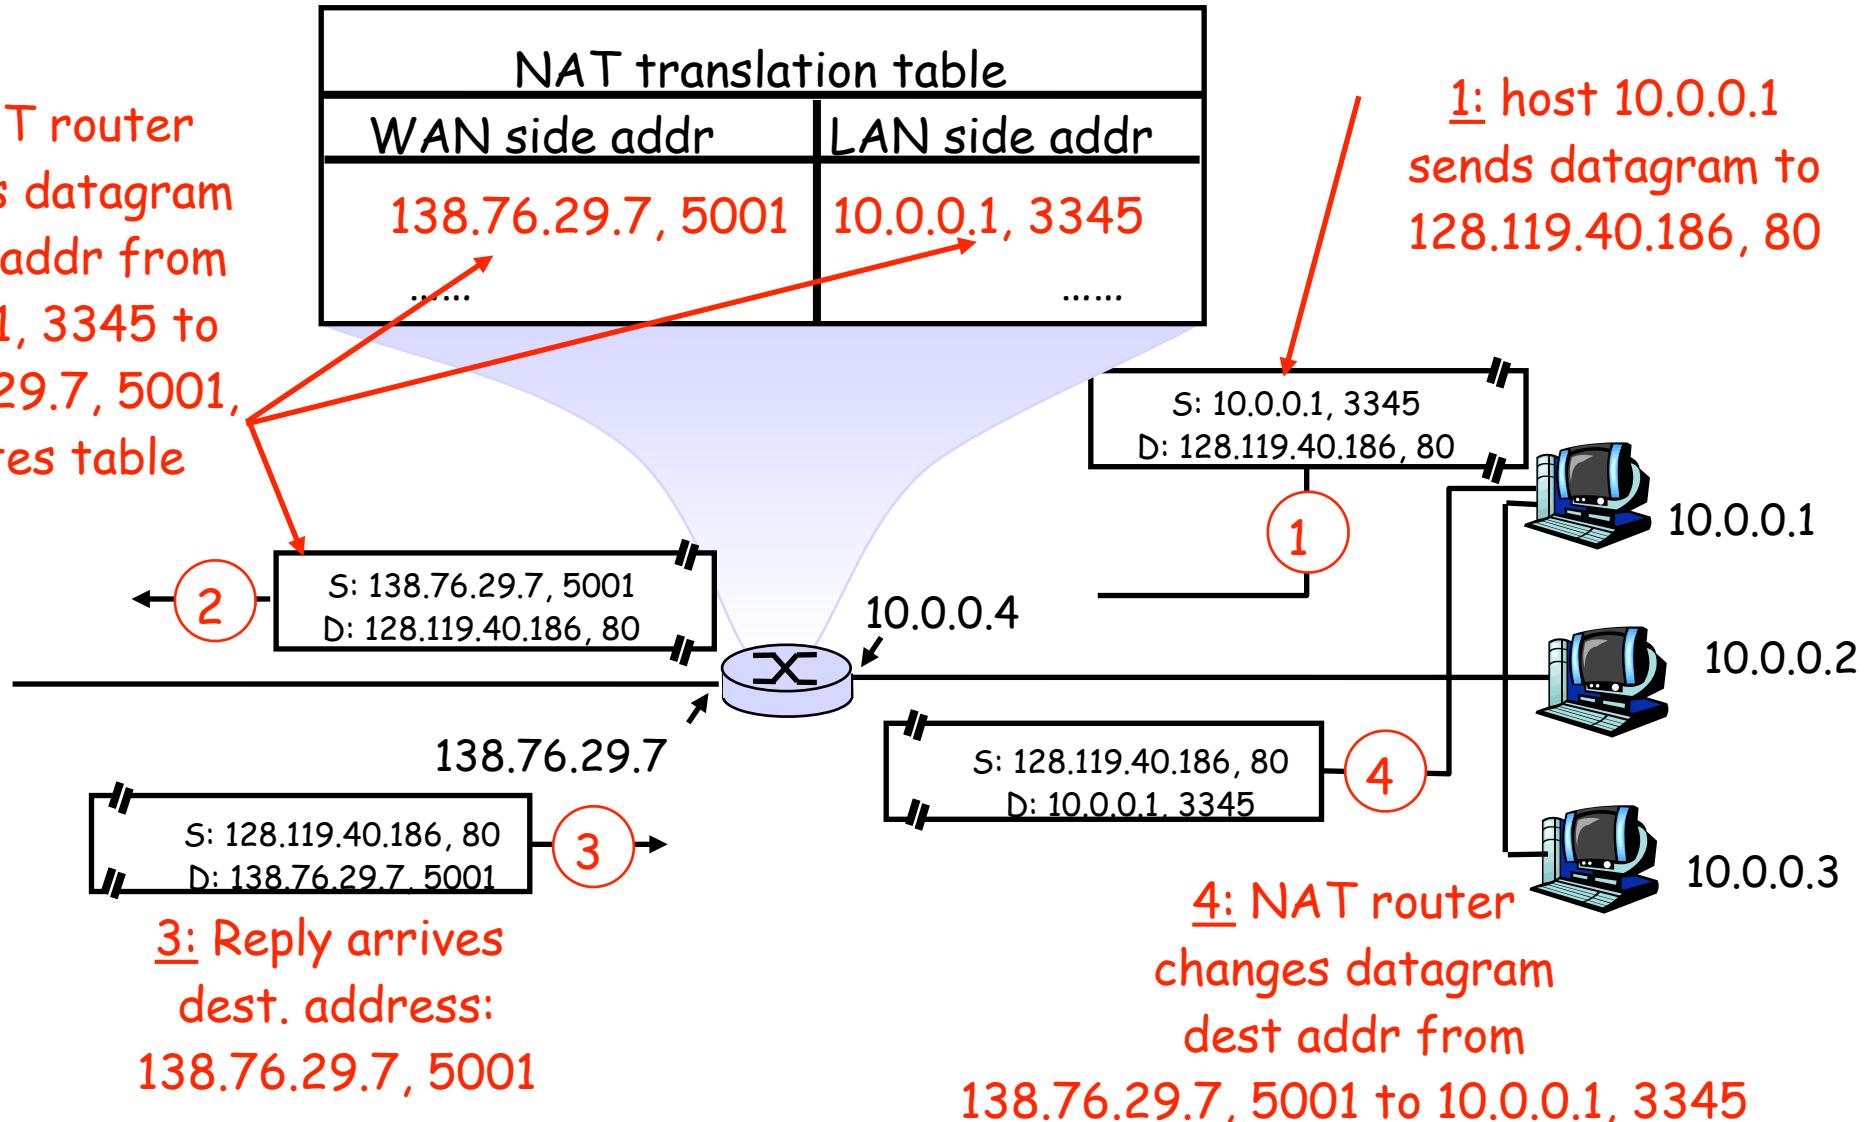
NAT: Network Address Translation

Implementation: NAT router must:

- ▶ **outgoing datagrams: replace** (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
... remote clients/servers will respond using (NAT IP address, new port #) as destination addr.
- ▶ **remember (in NAT translation table)** every (source IP address, port #) to (NAT IP address, new port #) translation pair
- ▶ **incoming datagrams: replace** (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: Network Address Translation

2: NAT router changes datagram source addr from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table

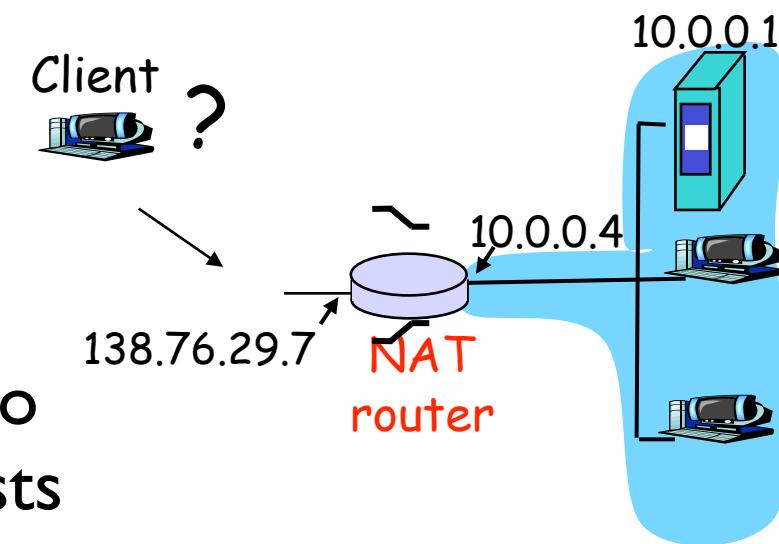


NAT: Network Address Translation

- 16-bit port-number field:
 - 60,000+ simultaneous connections with a single LAN-side address!
- NAT is controversial:
 - routers should only process up to layer 3
 - violates end-to-end argument
 - NAT possibility must be taken into account by app designers, eg, P2P applications
 - address shortage should instead be solved by IPv6
 - Anything idea why this might be a problem?

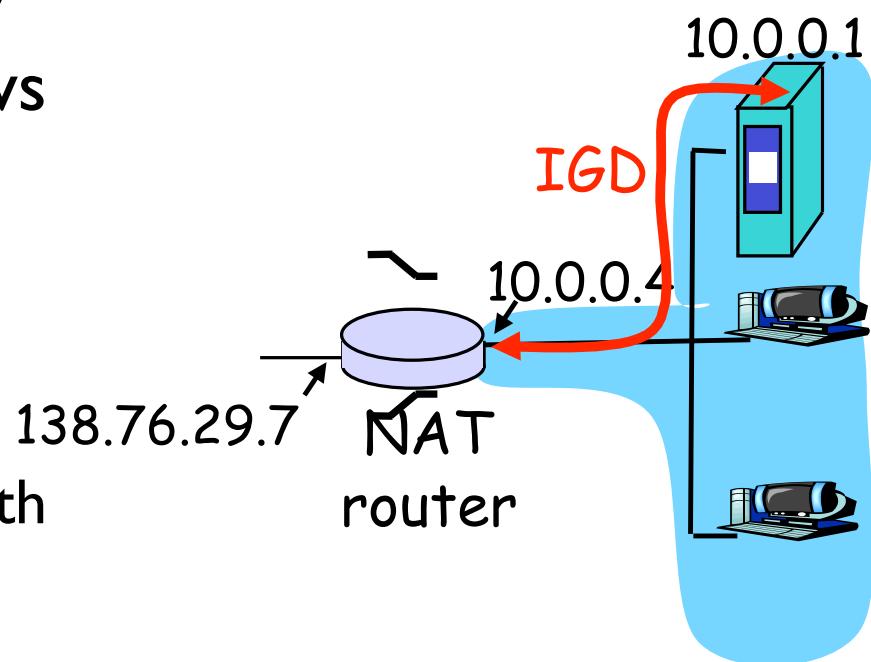
NAT traversal problem

- client wants to connect to server with address 10.0.0.1
 - ▶ server address 10.0.0.1 local to LAN (client can't use it as destination addr)
 - ▶ only one externally visible NATted address: 138.76.29.7
- solution I: statically configure NAT to forward incoming connection requests at given port to server
 - ▶ e.g., (123.76.29.7, port 2500) always forwarded to 10.0.0.1 port 25000



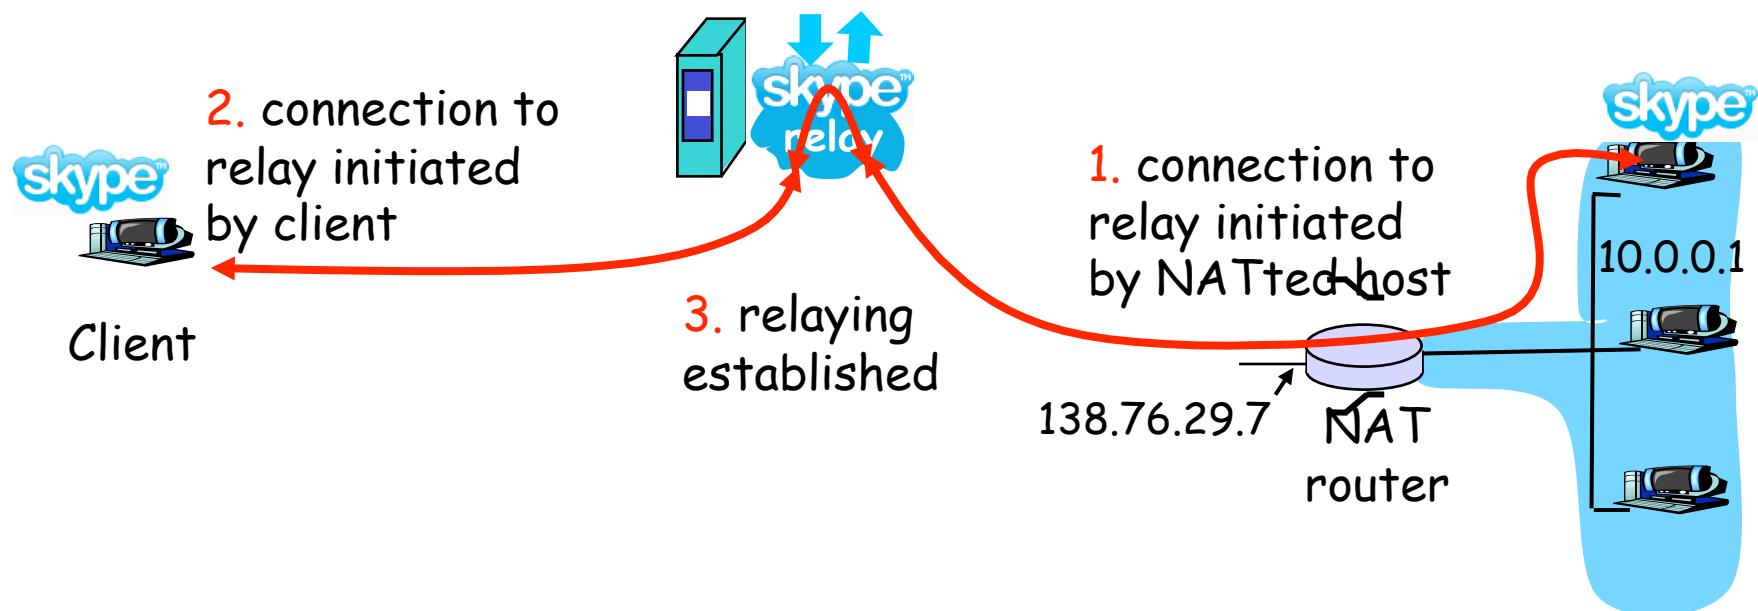
NAT traversal problem

- solution 2: Universal Plug and Play (UPnP) Internet Gateway Device (IGD) Protocol. Allows NATted host to:
 - ▶ learn public IP address (138.76.29.7)
 - ▶ add/remove port mappings (with lease times)
- i.e., automate static NAT port map configuration



NAT traversal problem

- solution 3: relaying (used in Skype)
 - NATed client establishes connection to relay
 - External client connects to relay
 - relay bridges packets between two connections



Chapter 4: Network Layer

- 4.1 Introduction
- 4.2 Virtual circuit and datagram networks
- 4.3 What's inside a router
- 4.4 IP: Internet Protocol
 - ▶ Datagram format
 - ▶ IPv4 addressing
 - ▶ ICMP
 - ▶ IPv6
- 4.5 Routing algorithms
 - ▶ Link state
 - ▶ Distance Vector
 - ▶ Hierarchical routing
- 4.6 Routing in the Internet
 - ▶ RIP
 - ▶ OSPF
 - ▶ BGP
- 4.7 Broadcast and multicast routing

ICMP: Internet Control Message Protocol

- used by hosts & routers to communicate network-level information
 - error reporting: unreachable host, network, port, protocol
 - echo request/reply (used by ping)
- network-layer “above” IP:
 - ICMP msgs carried in IP datagrams
- **ICMP message:** type, code plus first 8 bytes of IP datagram causing error

Type	Code	description
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header

Traceroute and ICMP

- Source sends series of UDP segments to dest
 - First has TTL = 1
 - Second has TTL=2, etc.
 - Unlikely port number
- When nth datagram arrives to nth router:
 - Router discards datagram
 - And sends to source an ICMP message (type 11, code 0)
 - Message includes name of router& IP address
- When ICMP message arrives, source calculates RTT
- Traceroute does this 3 times

Stopping criterion

- UDP segment eventually arrives at destination host
- Destination returns ICMP “port unreachable” packet (type 3, code 3)
- When source gets this ICMP, stops.

Smurf Attack

- ICMP Messages can be used in a classic “amplification” attack.
- An ICMP “ping” is sent to the broadcast address in a subnet (255.255.255.255) or network (192.168.1.255).
- All hosts receiving this message would automatically respond, thereby clogging the network.
 - ▶ Only took one message to initiate.



Chapter 4: Network Layer

- 4.1 Introduction
- 4.2 Virtual circuit and datagram networks
- 4.3 What's inside a router
- 4.4 IP: Internet Protocol
 - ▶ Datagram format
 - ▶ IPv4 addressing
 - ▶ ICMP
 - ▶ IPv6
- 4.5 Routing algorithms
 - ▶ Link state
 - ▶ Distance Vector
 - ▶ Hierarchical routing
- 4.6 Routing in the Internet
 - ▶ RIP
 - ▶ OSPF
 - ▶ BGP
- 4.7 Broadcast and multicast routing

IPv6

- **Initial motivation:** 32-bit address space completely allocated as of Feb. 2011.
- Additional motivation:
 - ▶ header format helps speed processing/forwarding
 - ▶ header changes to facilitate QoS

IPv6 datagram format:

- ▶ fixed-length 40 byte header
- ▶ no fragmentation allowed

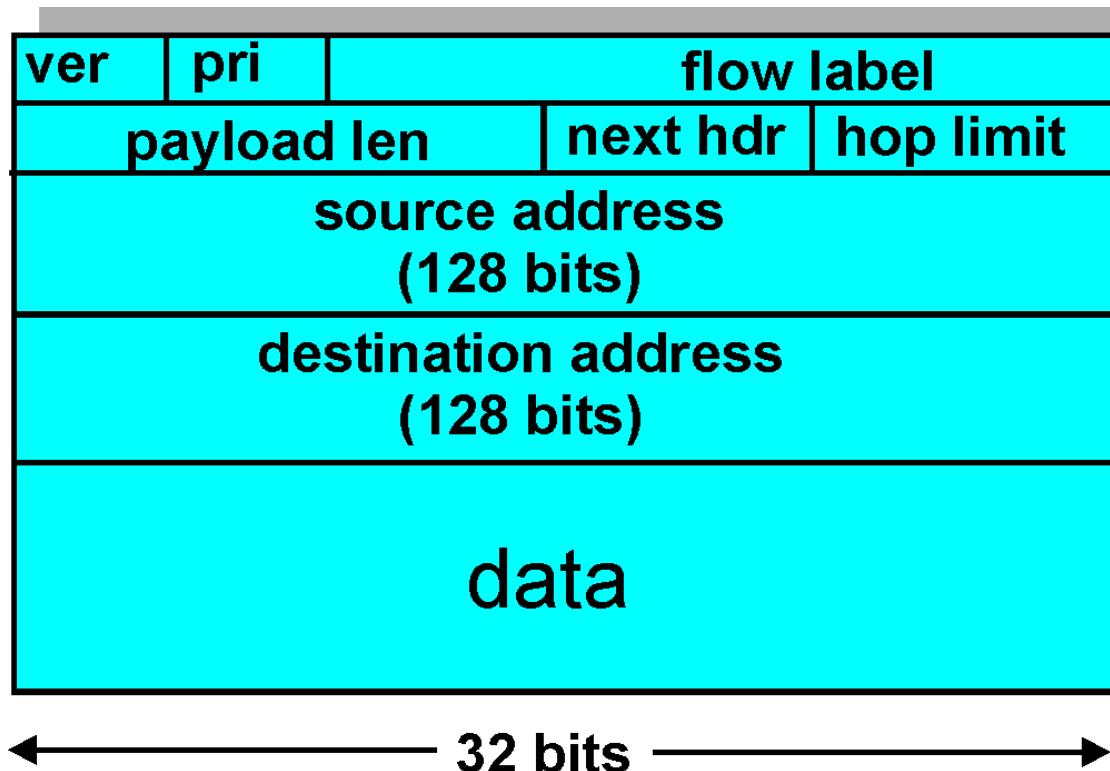
IPv6 Header (Cont)

Priority: identify priority among datagrams in flow

Flow Label: identify datagrams in same “flow.”

(concept of “flow” not well defined).

Next header: identify upper layer protocol for data



Other Changes from IPv4

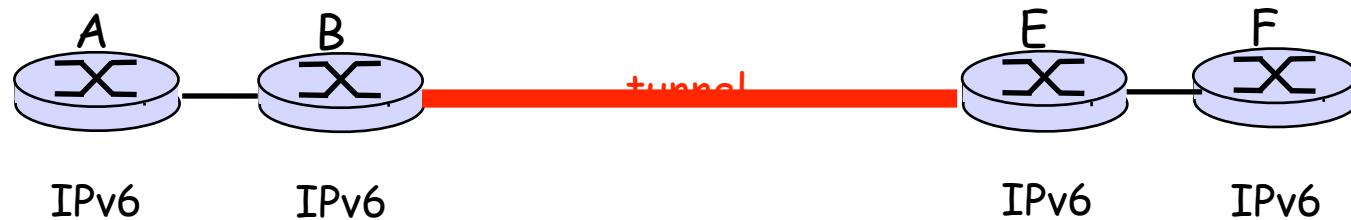
- **Checksum:** removed entirely to reduce processing time at each hop
- **Options:** allowed, but outside of header; indicated by “Next Header” field
- **ICMPv6:** new version of ICMP
 - ▶ additional message types, e.g. “Packet Too Big”
 - ▶ multicast group management functions

Transition From IPv4 To IPv6

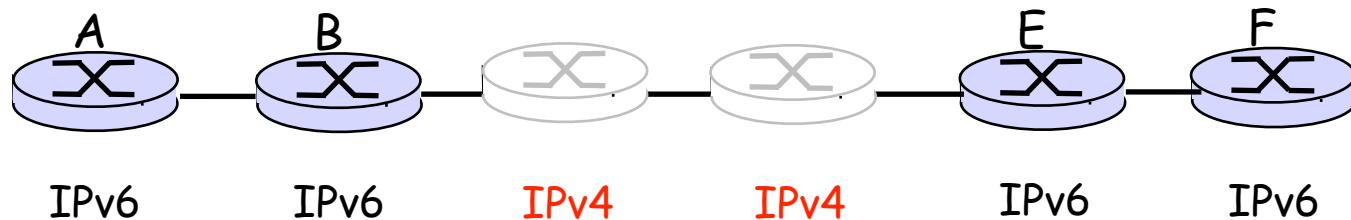
- Not all routers can be upgraded simultaneous
 - ▶ no “flag days”
 - ▶ How will the network operate with mixed IPv4 and IPv6 routers?
- **Tunneling:** IPv6 carried as payload in IPv4 datagram among IPv4 routers

Tunneling

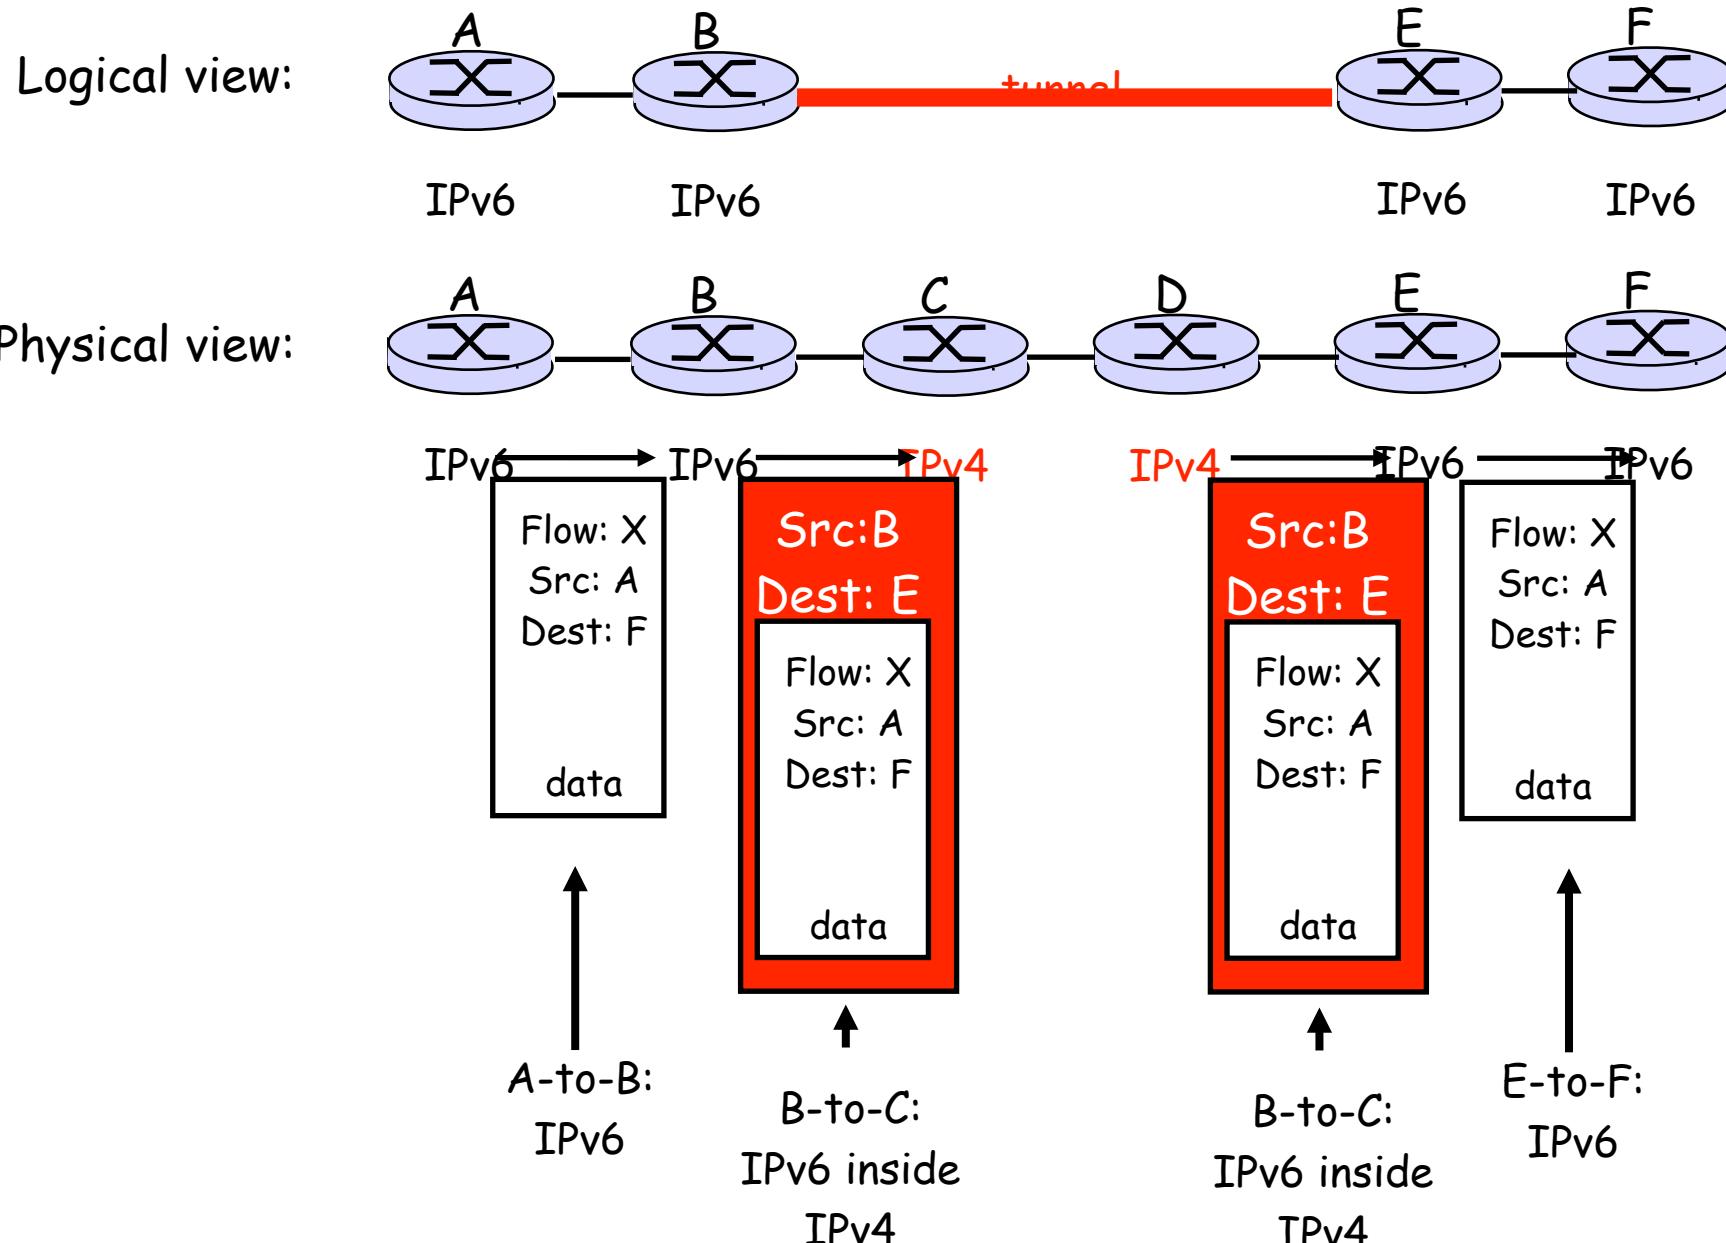
Logical view:



Physical view:



Tunneling



Next Time

- Read Section 4.5
 - ▶ Routing algorithms - this is important stuff

